
Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network

Christian Ledig et al

CVPR 2017

Presenter



20212047, Yechan Kim | yechankim@gm.gist.ac.kr

Please read ‘README.md’ before scoring.

The screenshot shows a text editor window with the title "README.md". The content is as follows:

SRGAN

Unofficial Implementation for "Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Model"

Yechan Kim (20212047)

Prerequisites

- See `requirements.txt` for details.

```
torch >= 1.10.0
torchvision >= 0.11.1
glob2
numpy
matplotlib
tqdm
tensorboard
tensorboardx
```

Toy-Examples

- This repository provides DIV2K (with only few examples) and Set5 for convenience. Note that given DIV2K has only few examples, thus you need to download the full DIV2K dataset for training.
- You can download DIV2K or other datasets for super-resolution task via internet.

...

(Recap) Super-resolution?



*Super-
resolution*



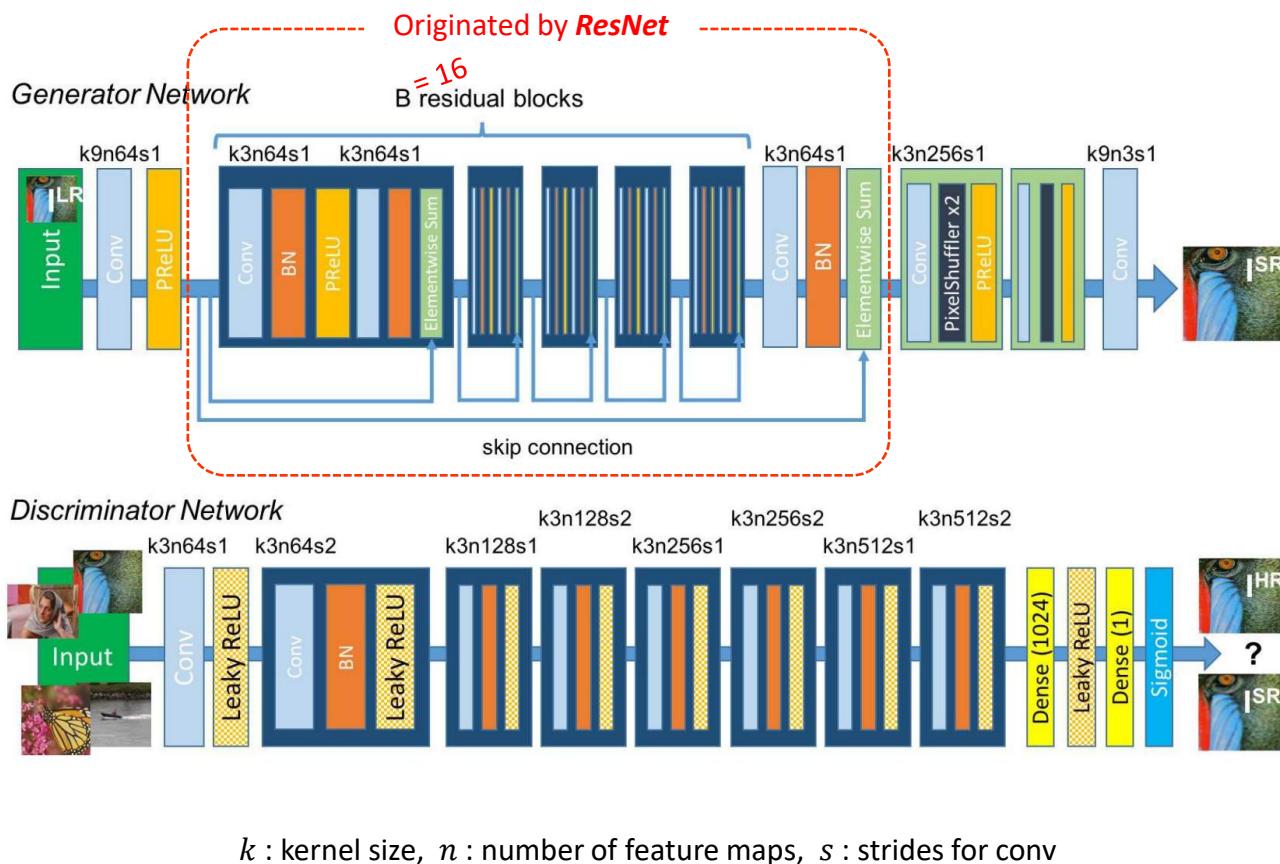
- Super-resolution is the task of taking an input of a low resolution (LR) and **upsampling it to that of a high resolution**.
- The **ill-posed nature of the underdetermined SR problem** is particularly pronounced for high upscaling factors, for which **texture detail** in the reconstructed SR images is typically absent [1].

Reference:

[1] Ledig, Christian, et al. "Photo-realistic single image super-resolution using a generative adversarial network." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017.

(Recap) SRGAN?

- Architectures



- Avoid using any pooling layers.
- Generative network (**SRResNet**) is pre-trained alone first before adversarial training.
- Then, this work alternatively trains the discriminator and generator. → *Details on the next slide!*

(Recap) SRGAN? [Code]

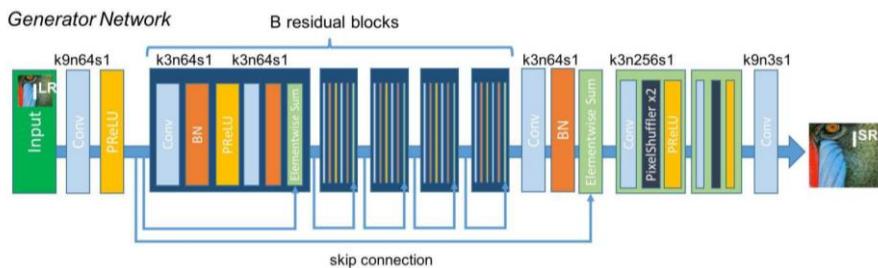
- SRGAN/Model/Generator.py

44 lines (41 sloc) | 2.03 KB

```

1  from torch import nn
2  from Model.Block import ResidualBlock, list_to_sequential
3  from Model.Block import UpsamplingBlock
4  import numpy as np
5
6
7  class Generator(nn.Module):
8      def __init__(self, in_channels=3, num_residual_blocks=16):
9          super(Generator, self).__init__() # (in, out)
10         num_upsampling_blocks = 2
11         out_channels = 64
12         self.input_block = nn.Sequential(
13             nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=(9, 9), # (3, 64)
14                     stride=1, padding=4, bias=False),
15             nn.PReLU())
16
17         self.residual_blocks = list_to_sequential(
18             [ResidualBlock(in_channels=out_channels, out_channels=out_channels) # (64, 64)
19              for _ in range(num_residual_blocks)])
20
21         self.intermediate_block = nn.Sequential(
22             nn.Conv2d(in_channels=out_channels, out_channels=out_channels, kernel_size=(3, 3), # (64, 64)
23                     stride=1, padding=1, bias=False),
24             nn.BatchNorm2d(num_features=out_channels))
25
26         # for connecting residual blocks with upsampling blocks!
27         self.upsampling_blocks = list_to_sequential(
28             [UpsamplingBlock(in_channels=out_channels, # (64, 64 * num_upsampling_blocks ^ 2)
29                             scaling_factor=num_upsampling_blocks)
30                 for _ in range(num_upsampling_blocks)])
31
32         self.output_block = nn.Sequential(
33             nn.Conv2d(in_channels=out_channels, out_channels=in_channels, kernel_size=(9, 9),
34                     stride=1, padding=4, bias=False),
35             # we might add extra non-linear activation here!
36         )
37         self.short_cut = nn.Sequential()
38
39     def forward(self, x):
40         identity = self.input_block(x)
41         out = self.residual_blocks(identity)
42         out = self.short_cut(identity) + self.intermediate_block(out)
43         out = self.upsampling_blocks(out)
44         out = self.output_block(out)
45
46         return out

```



SRGAN/Model/Block.py

56 lines (45 sloc) | 2.11 KB

```

1  from torch import nn
2
3
4  class ResidualBlock(nn.Module): # for Generator (1)
5      def __init__(self, in_channels, out_channels, kernel_size=3, stride=1, padding=1):
6          super(ResidualBlock, self).__init__()
7          self.conv1 = nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=kernel_size,
8                               stride=stride, padding=padding, bias=False)
9          self.conv2 = nn.Conv2d(in_channels=out_channels, out_channels=out_channels, kernel_size=kernel_size,
10                           stride=stride, padding=padding, bias=False)
11         self.bn = nn.BatchNorm2d(num_features=out_channels)
12         self.p_relu = nn.PReLU()
13         self.short_cut = nn.Sequential()
14
15     def forward(self, x):
16         identity = x
17         out = self.conv1(x)
18         out = self.bn(out)
19         out = self.p_relu(out)
20         out = self.conv2(out)
21         out = self.bn(out)
22         out = self.short_cut(identity) + out
23
24         return out
25
26 class UpsamplingBlock(nn.Module): # for Generator (2)
27     def __init__(self, in_channels, scaling_factor, kernel_size=(3, 3), stride=1, padding=1):
28         super(UpsamplingBlock, self).__init__()
29         out_channels = in_channels * (scaling_factor ** 2) # -> Why?
30         self.conv = nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=kernel_size,
31                             stride=stride, padding=padding, bias=False)
32         self.pixel_shuffle = nn.PixelShuffle(scaling_factor)
33         self.p_relu = nn.PReLU()
34
35     def forward(self, x):
36         out = self.conv(x)
37         out = self.pixel_shuffle(out)
38         out = self.p_relu(out)
39
40         return out
41
42     ...
43
44     def list_to_sequential(layer_list):
45         return nn.Sequential(*layer_list)

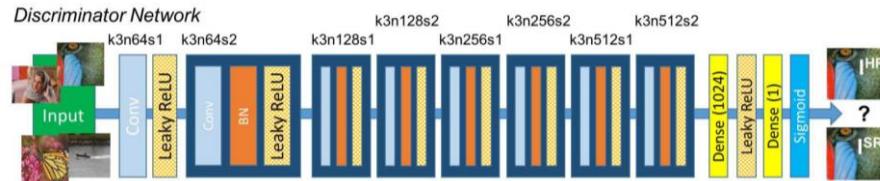
```

(Recap) SRGAN? [Code]

- SRGAN/Model/Discriminator.py

33 lines (30 sloc) | 1.28 KB

```
1 from torch import nn
2 from Model.Block import ConvBlock, list_to_sequential
3
4
5 class Discriminator(nn.Module):
6     def __init__(self, in_channels=3, conv_in_out_chs_strides=None):
7         super(Discriminator, self).__init__()
8         self.linear_width = 1024
9         if not conv_in_out_chs_strides:
10             conv_in_out_chs_strides = [(64, 64, 2), (64, 128, 1), (128, 128, 2),
11                                         (128, 256, 1), (256, 512, 1), (512, 512, 2)]
12         self.input_block = nn.Sequential(
13             nn.Conv2d(in_channels=in_channels, out_channels=64, kernel_size=(3, 3),
14                       stride=1, padding=1, bias=False),
15             nn.LeakyReLU(0.2)
16         )
17         self.conv_blocks = list_to_sequential(
18             [ConvBlock(in_channels=in_ch, out_channels=out_ch, stride=stride)
19              for in_ch, out_ch, stride in conv_in_out_chs_strides]
20         )
21         self.output_block = nn.Sequential(
22             nn.Linear(out_features=self.linear_width),
23             nn.LeakyReLU(0.2),
24             nn.AdaptiveAvgPool2d(1),
25             nn.Sigmoid()
26         )
27
28     def forward(self, x):
29         out = self.input_block(x)
30         out = self.conv_blocks(out)
31         out = out.view(-1, self.linear_width)
32         out = self.output_block(out)
33         return out
```



SRGAN/Model/Block.py

56 lines (45 sloc) | 2.11 KB

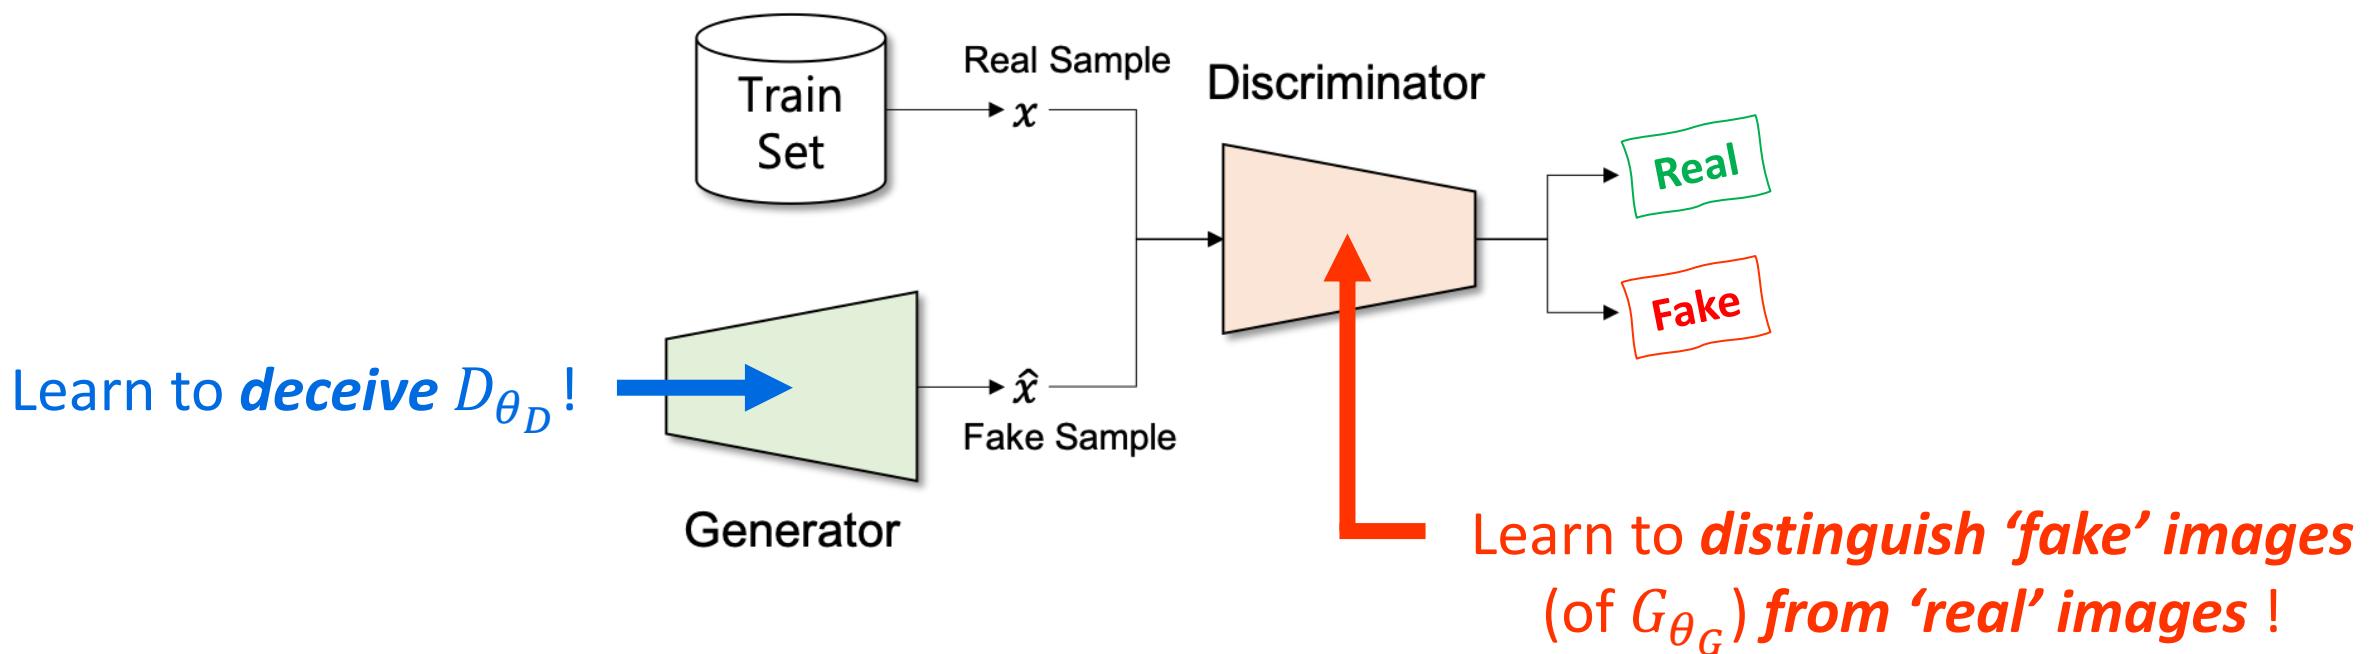
```
1 from torch import nn
...
42     class ConvBlock(nn.Module): # for Discriminator
43         def __init__(self, in_channels, out_channels, kernel_size=3, stride=1, padding=1):
44             super(ConvBlock, self).__init__()
45             self.block = nn.Sequential(
46                 nn.Conv2d(in_channels, out_channels, kernel_size, stride, padding),
47                 nn.BatchNorm2d(out_channels),
48                 nn.LeakyReLU(0.2)
49             )
50
51         def forward(self, x):
52             return self.block(x)
53
54
55     def list_to_sequential(layer_list):
56         return nn.Sequential(*layer_list)
```

(Recap) SRGAN?

- Training Objective

- Let I^{LR} and I^{HR} be images with low resolution and high resolution, respectively.
- Let G_{θ_G} and D_{θ_D} be the generative and discriminator networks, respectively.
- The optimization goal is to solve the adversarial min-max problem as follows:

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))].$$



(Recap) SRGAN?

- Training Objective (cont.)

Discriminator

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$$



$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} \underbrace{[\log D_{\theta_D}(I^{HR})]}_{\text{Maximum when } D_{\theta_D}(I^{HR}) = 1} + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$$

Maximum when $D_{\theta_D}(I^{HR}) = 1$

(\because Range of $D_{\theta_D}(\cdot)$: [0, 1].)

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} \underbrace{[\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]}_{\text{Maximum when } D_{\theta_D}(G_{\theta_G}(I^{LR})) = 0}$$

Maximum when $D_{\theta_D}(G_{\theta_G}(I^{LR})) = 0$
 $(\because$ Range of $D_{\theta_D}(\cdot)$: [0, 1].)

(Recap) SRGAN?

- Training Objective (cont.)

Generator

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$$



~~$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$$~~

→ G_{θ_G} is independent of this term.

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\underbrace{\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))}_{\text{Minimum when } D_{\theta_D}(G_{\theta_G}(I^{LR})) = 1}]$$

Minimum when $D_{\theta_D}(G_{\theta_G}(I^{LR})) = 1$
 $(\because \log 0 = -\infty.)$

Discriminator

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$$

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\underbrace{\log D_{\theta_D}(I^{HR})}_{\text{Maximum when } D_{\theta_D}(I^{HR}) = 1}] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))]$$

Maximum when $D_{\theta_D}(I^{HR}) = 1$
 $(\because \text{Range of } D_{\theta_D}(\cdot): [0, 1].)$

$$\min_{\theta_G} \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\underbrace{\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))}_{\text{Maximum when } D_{\theta_D}(G_{\theta_G}(I^{LR})) = 0}]$$

Maximum when $D_{\theta_D}(G_{\theta_G}(I^{LR})) = 0$
 $(\because \text{Range of } D_{\theta_D}(\cdot): [0, 1].)$

(Recap) Perceptual Loss?

- [1] adopts a **perceptual loss** function for training the generator as follows:

$$l^{SR} = l_X^{SR} + 10^{-3} l_{gen}^{SR}.$$

- l_X^{SR} : Content Loss

- This work uses $l_{VGG / i,j}^{SR} = \frac{1}{W_{i,j}H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$,

where $W_{i,j}$ and $H_{i,j}$ are the dimensions of the respective feature map $\phi_{i,j}$

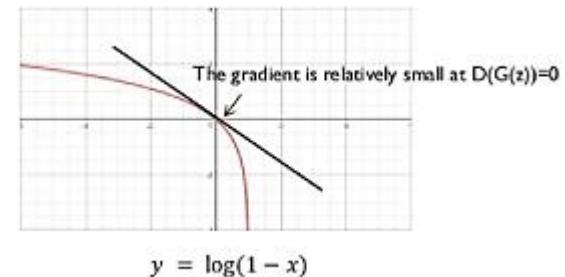
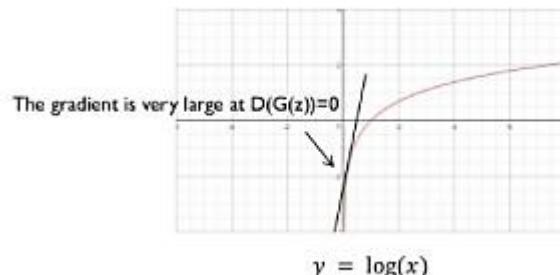
which is obtained by *j-th convolution* (after activation) *before i-th max-pooling* layer within VGG-19.

- l_{gen}^{SR} : Adversarial Loss

- $I_{Gen}^{SR} = \sum_{n=1}^N -\log D_{\theta_D}(G_{\theta_G}(I^{LR}))$



$$I_{Gen}^{SR} = \sum_{n=1}^N \log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))$$



Reference:

[1] Ledig, Christian, et al. "Photo-realistic single image super-resolution using a generative adversarial network." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017.

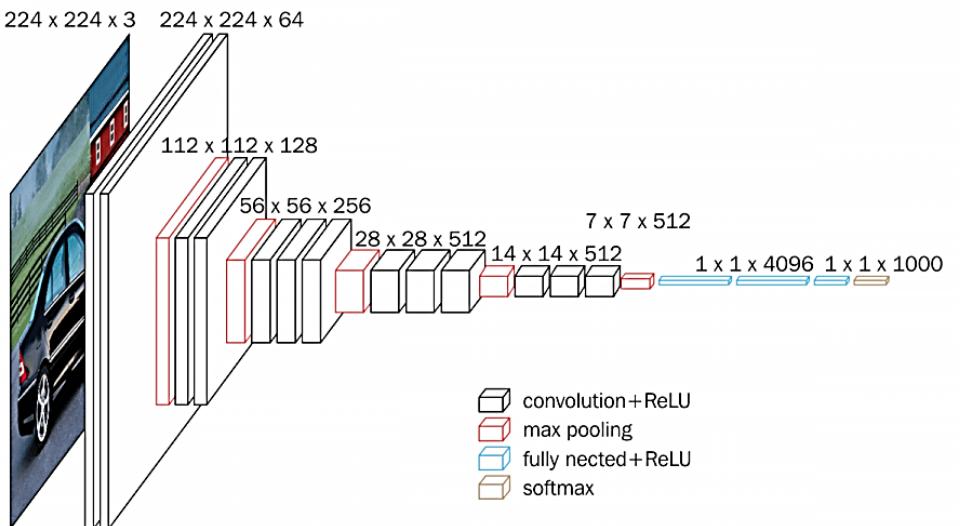
[3] <https://www.slideshare.net/NaverEngineering/1-gangenerative-adversarial-network>

(+ α) How to Enhance the Content Loss?

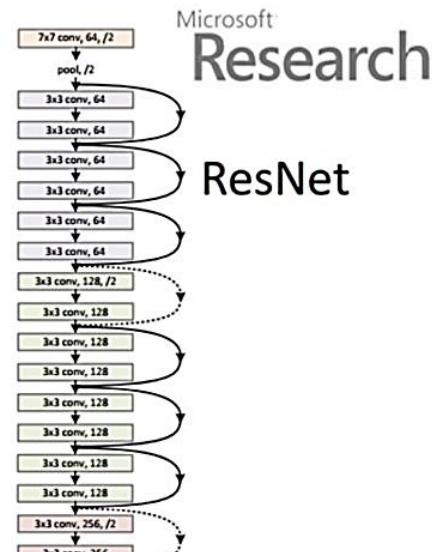
- l_X^{SR} : Content Loss

$$\bullet \quad l_{VGG}^{SR} / i,j = \frac{1}{W_{i,j}H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} \left(\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}\left(G_{\theta_G}(I^{LR})\right)_{x,y} \right)^2.$$

↳ How about designing the ‘content loss’ based on *ResNet* [4]?



VGGNet



ResNet

Reference:

- [1] Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." arXiv preprint arXiv:1409.1556 (2014).
[4] He, Kaiming, et al. "Deep residual learning for image recognition." Proceedings of the IEEE conference on computer vision and pattern recognition. 2016.

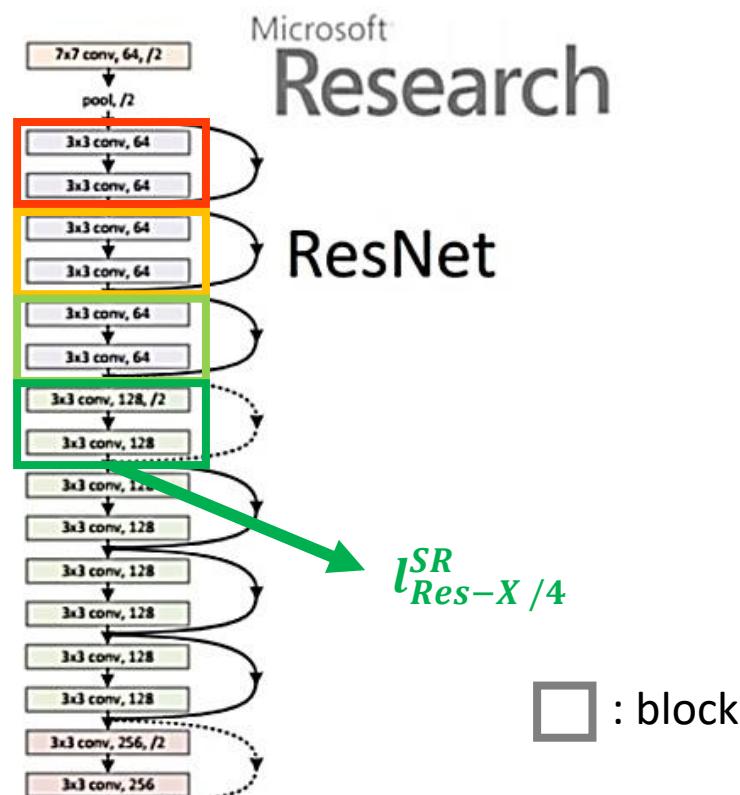
$$I_{\text{VGG}}^{SR} / i,j = \frac{1}{W_{i,j}H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} \left(\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}\left(G_{\theta_G}(I^{LR})\right)_{x,y} \right)^2$$

(+α) How to Enhance the Content Loss? (cont.)

- l_X^{SR} : Content Loss (**Proposed**)

$$\bullet \quad l_{\text{Res}-X / i}^{SR} = \frac{1}{W_i H_i} \sum_{x=1}^{W_i} \sum_{y=1}^{H_i} \left(\phi_i(I^{HR})_{x,y} - \phi_i\left(G_{\theta_G}(I^{LR})\right)_{x,y} \right)^2,$$

where W_i and H_i are the dimensions of the respective feature map ϕ_i which is obtained by ' i -th res-block' within ResNet-X.



Datasets

- For training the SRResNet and SRGAN models,

DIV2K

Introduced by Eirikur Agustsson et al. in [NTIRE 2017 Challenge on Single Image Super-Resolution: Dataset and Study](#)

DIV2K is a popular single-image super-resolution dataset which contains 1,000 images with different scenes and is split into 800 for training, 100 for validation and 100 for testing. It was collected for NTIRE2017 and NTIRE2018 Super-Resolution Challenges in order to encourage research on image super-resolution with more realistic degradation. This dataset contains low resolution images with different types of degradations. Apart from the standard bicubic downsampling, several types of degradations are considered in synthesizing low resolution images for different tracks of the challenges. Track 2 of NTIRE 2017 contains low resolution images with unknown x4 downscaling. Track 2 and track 4 of NTIRE 2018 correspond to realistic mild x4 and realistic wild x4 adverse conditions, respectively. Low-resolution images under realistic mild x4 setting suffer from motion blur, Poisson noise and pixel shifting. Degradations under realistic wild x4 setting are further extended to be of different levels from image to image.

Source: [Unsupervised Image Super-Resolution with an Indirect Supervised Path](#)

[Homepage](#)

Benchmarks

[Edit](#)



Usage ▲



- Train:
- Validation:

[DIV2K_train_HR/*.png](#)
[DIV2K_valid_HR/*.png](#)

Trend	Task	Dataset Variant	Best Model	Paper	Code
	Image Super-Resolution	DIV2K val - 16x upscaling	ABPN	Paper	Code
	Image Super-Resolution	DIV2K val - 2x upscaling	CAR	Paper	Code
	Image Super-Resolution	DIV2K val - 4x upscaling	CAR	Paper	Code

[Edit](#)

[Custom \(research-only\)](#)

Datasets

- For evaluating the SRResNet and SRGAN models,

Set5

Introduced by Marco Bevilacqua et al. in [Low-Complexity Single-Image Super-Resolution based on Nonnegative Neighbor Embedding](#)

The Set5 dataset is a dataset consisting of 5 images ("baby", "bird", "butterfly", "head", "woman") commonly used for testing performance of Image Super-Resolution models.

[Homepage](#)

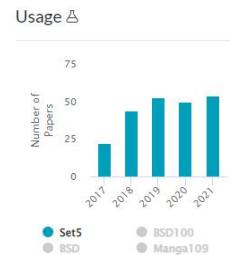
Benchmarks

Trend	Task	Dataset Variant	Best Model	Paper	Code
	Image Super-Resolution	Set5 - 4x upscaling	SwinIR	Paper	Code
	Image Super-Resolution	Set5 - 2x upscaling	CAR	Paper	Code
	Image Super-Resolution	Set5 - 3x upscaling	DRLN+	Paper	Code
	Image Super-Resolution	Set5 - 8x upscaling	MFSRCNN	Paper	Code
	Compressive Sensing	Set5	LapCSNet	Paper	Code

[Edit](#)



Source: <http://people.rennes.inria.fr/Aline.Ro...>



Set14

Introduced by Roman Zeyde et al. in [On Single Image Scale-Up Using Sparse-Representations](#)

The Set14 dataset is a dataset consisting of 14 images commonly used for testing performance of Image Super-Resolution models.

[Homepage](#)

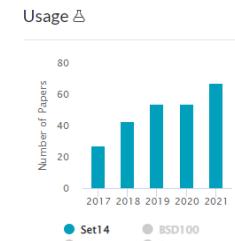
Benchmarks

Trend	Task	Dataset Variant	Best Model	Paper	Code
	Image Super-Resolution	Set14 - 4x upscaling	SwinIR	Paper	Code
	Image Super-Resolution	Set14 - 2x upscaling	CAR	Paper	Code
	Image Super-Resolution	Set14 - 3x upscaling	DRLN+	Paper	Code
	Image Super-Resolution	Set14 - 8x upscaling	DBPN-RES-MR64-3	Paper	Code

[Edit](#)



Source: <https://www.ece.rice.edu/~wakin/im...>



- HR: [Set5/GTmod12/*.png](#)
- LR: [Set5/LRbicx4/*.png](#)

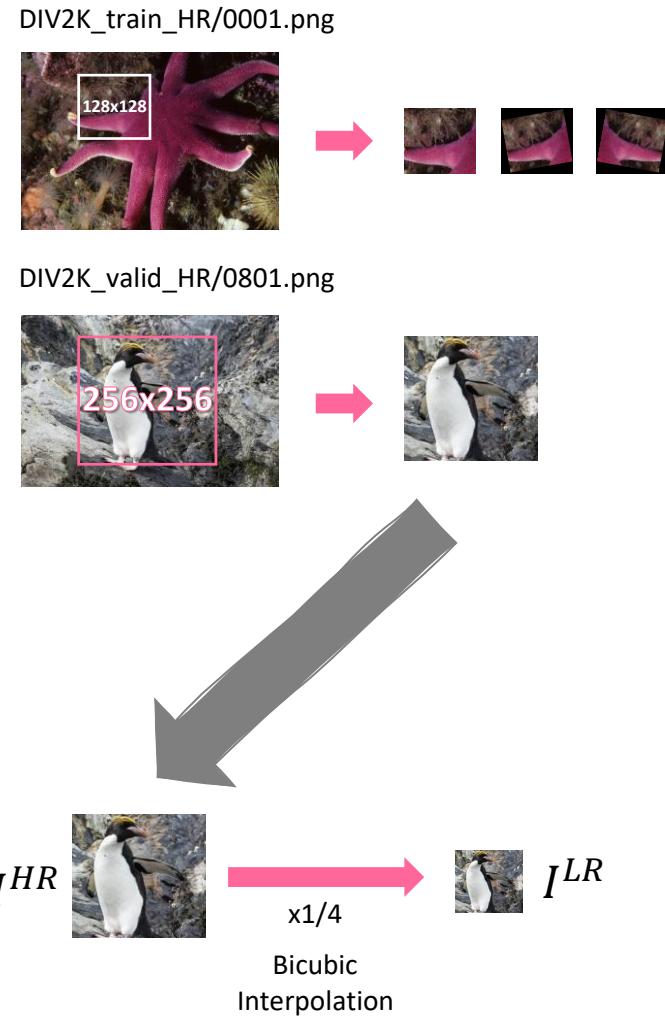
- HR: [Set14/GTmod12/*.png](#)
- LR: [Set14/LRbicx4/*.png](#)

Reference:

- [6] <https://paperswithcode.com/dataset/set5>
- [7] <https://paperswithcode.com/dataset/set11>

How to Load HR and LR Images?

- HR (High Resolution) images
 - This work only downloads **DIV2K_train_HR** and **DIV2K_valid_HR**.
 - Zero-padding, Random-cropping, Random-rotation, and horizontal-flipping are applied to training set of DIV2K.
 - 128x128 random-cropping
 - 10-degree random-rotation
 - Horizontal flipping with probability 0.5
 - Only center-cropping is applied to validation set of DIV2K.
 - 256x256 center-cropping
- LR (Low Resolution) images
 - LR images are generated via resizing pre-processed HR images by a quarter (1/4) for both training and validation.
 - Bicubic interpolation is used for reducing the size of images.



How to Load HR and LR Images?

[Code]

SRGAN/Dataset/dataset.py

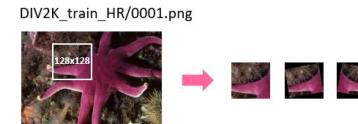
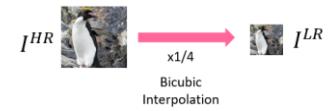
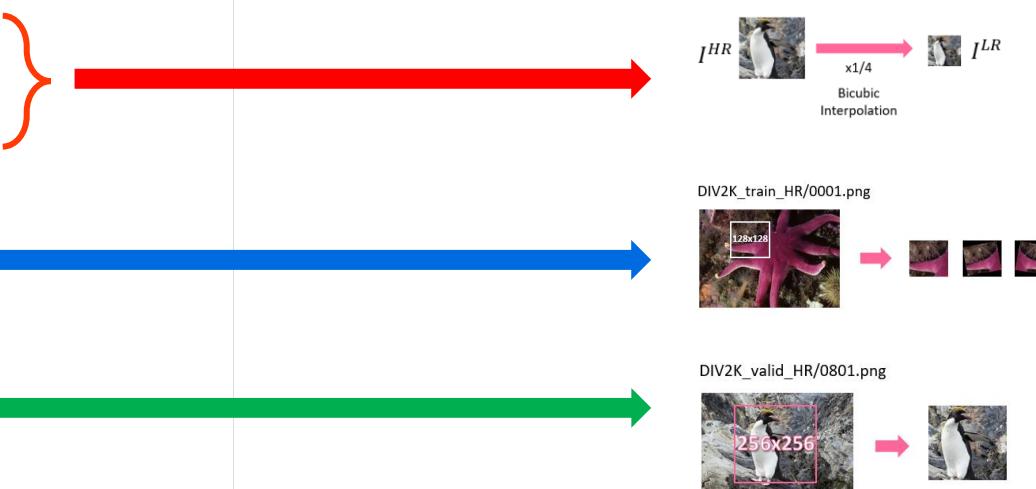
99 lines (88 sloc) | 3.71 KB

Raw Blame

```
1 from PIL import Image
2
3 import glob
4 from torch.utils.data import Dataset
5 from torchvision.transforms import transforms, InterpolationMode
6
7
8 class HR_to_LR_PairDataset(Dataset):
9     def __init__(self, data_path, mode, crop_size=128, rescale_ratio=4):
10         self.img_files = sorted(glob.glob(f'{data_path}/*.*'))
11         assert len(self.img_files) > 0, "No images exist in the given dataset path."
12         self.LR_transforms = self.get_LR_transforms(crop_size, rescale_ratio)
13         self.HR_transforms = self.get_HR_transforms(mode, crop_size)
14
15     def get_LR_transforms(self, crop_size, rescale_ratio):
16         return transforms.Compose([
17             transforms.ToPILImage(),
18             transforms.Resize(crop_size // rescale_ratio, interpolation=InterpolationMode.BICUBIC),
19             transforms.ToTensor()
20         ])
21
22     def get_HR_transforms(self, mode, crop_size):
23         if mode == 'train':
24             return transforms.Compose([
25                 transforms.RandomCrop(crop_size),
26                 transforms.RandomRotation(10),
27                 transforms.RandomHorizontalFlip(0.5),
28                 transforms.ToTensor()
29             ])
30         else:
31             return transforms.Compose([
32                 transforms.CenterCrop(crop_size),
33                 transforms.ToTensor()
34             ])
35
36     def __len__(self):
37         return len(self.img_files)
38
39     def __getitem__(self, idx):
40         hr_img = self.HR_transforms(Image.open(self.img_files[idx]))
41         lr_img = self.LR_transforms(hr_img)
42         return lr_img, hr_img
```

The code block contains three annotations:

- A red curly brace on the right side of the code, spanning from the end of the `__init__` method to the end of the file, with a red horizontal bar extending to its right.
- A blue curly brace on the right side of the `get_LR_transforms` method, spanning from its opening brace to its closing brace.
- A green curly brace on the right side of the `get_HR_transforms` method, spanning from its opening brace to its closing brace.



Training Details for SRResNet

- SRResNet models are trained for 250 epochs on DIV2K.
- During training, the learning rate is 0.0001 and no learning rate scheduler is adopted.
- ADAM optimizer is used with betas=(0.9, 0.999).
- This work considers various training objectives as follows:
 - [Default] **MSE** Loss

$$l_{\text{MSE}}^{SR} = \frac{1}{WH} \sum_{x=1}^W \sum_{y=1}^H \left((I^{HR})_{x,y} - \left(G_{\theta_G}(I^{LR}) \right)_{x,y} \right)^2,$$

where W and H are the dimensions of the images.



- [+a] Content Loss – **VGG** Loss

$$l_{\text{VGG}}^{SR} / i,j = \frac{1}{W_{i,j} H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} \left(\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}\left(G_{\theta_G}(I^{LR})\right)_{x,y} \right)^2,$$

where $W_{i,j}$ and $H_{i,j}$ are the dimensions of the respective feature map $\phi_{i,j}$

which is obtained by j -th convolution (after activation) before i -th max-pooling layer within VGG-19.



- [+a] Content Loss – **ResNet** Loss

$$l_{\text{Res}-X / i}^{SR} = \frac{1}{W_i H_i} \sum_{x=1}^{W_i} \sum_{y=1}^{H_i} \left(\phi_i(I^{HR})_{x,y} - \phi_i\left(G_{\theta_G}(I^{LR})\right)_{x,y} \right)^2,$$

where W_i and H_i are the dimensions of the respective feature map ϕ_i which is obtained by ' i -th res-block' within ResNet-X.



Training Details for SRResNet

- [+a] Content Loss – VGG Loss [Code]

$$l_{\text{VGG}}^{SR} / i,j = \frac{1}{W_{i,j} H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} \left(\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}\left(G_{\theta_G}(I^{LR})\right)_{x,y} \right)^2,$$

where $W_{i,j}$ and $H_{i,j}$ are the dimensions of the respective feature map $\phi_{i,j}$

which is obtained by j -th convolution (after activation) before i -th max-pooling layer within VGG-19.



SRGAN/ContentLoss/VGGLoss.py

42 lines (37 sloc) | 1.62 KB

Raw Blame

```
1 import torch
2 import torch.nn as nn
3
4
5 def get_extractor_for_VGG(torch_model, i=5, j=4):
6     """ [Logic]
7     extract the feature extractor from torch_model (vgg-model)
8     here, the extractor should include from 1-st layer to the 'j'-th convolutional layer
9         before 'i'-th max-pooling layer.
10    ...
11
12    layers = list(torch_model.children())[0]
13    last_conv_idx_list = []
14    last_maxpool_cnt = 0
15    for idx, layer in enumerate(layers):
16        layer.requires_grad = False # freezing!
17        if last_maxpool_cnt == i:
18            break
19        if 'MaxPool2d' in str(layer):
20            last_maxpool_cnt += 1
21        elif 'Conv2d' in str(layer):
22            last_conv_idx_list.append(idx + 1)
23    final_layers = layers[0:last_conv_idx_list[-(j + 1)]]
24    return nn.Sequential(*final_layers)
```

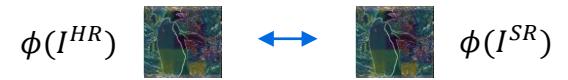
```
26 class VGGLoss(nn.Module):
27     def __init__(self, i=5, j=4, device='cuda:0', vgg_model_name='vgg19_bn'):
28         """ [Logic]
29         1) extract the feature map from the 'j'-th convolutional layer before 'i'-th max-pooling layer
30             for both hr_image (ground_truth) and sr_image (prediction).
31         2) compute and return mse_loss between features maps of hr_image & sr_image.
32         ...
33         super(VGGLoss, self).__init__()
34         vgg = torch.hub.load('pytorch/vision:v0.10.0', vgg_model_name, pretrained=True).to(device).eval()
35         self.extractor = get_extractor_for_VGG(vgg, i, j)
36         self.mse_loss = nn.MSELoss()
37
38     def forward(self, hr_img, sr_img):
39         hr_feature_map = self.extractor(hr_img)
40         sr_feature_map = self.extractor(sr_img)
41         vgg_loss = self.mse_loss(hr_feature_map, sr_feature_map)
42
43     return vgg_loss
```

Training Details for SRResNet

- [+a] Content Loss – ResNet Loss [Code]

$$l_{\text{Res}-X / i}^{SR} = \frac{1}{W_i H_i} \sum_{x=1}^{W_i} \sum_{y=1}^{H_i} \left(\phi_i(I^{HR})_{x,y} - \phi_i(G_{\theta_G}(I^{LR}))_{x,y} \right)^2,$$

where W_i and H_i are the dimensions of the respective feature map ϕ_i which is obtained by ‘ i -th res-block’ within ResNet-X.



SRGAN/ContentLoss/ResNetLoss.py

```
54 lines (48 sloc) | 1.79 KB
Raw Blame ⌂ ⌄ ⌁ ⌃ ⌅ ⌆
```

```
1 import torch
2 import torch.nn as nn
3
4
5 def get_extractor_for_ResNet(torch_model, i=6):
6     ''' [Logic]
7     extract the feature extractor from torch_model (ResNet-model)
8     here, the extractor should include from 1-st layer to 'i'-th residual block.
9     '''
10    layers = list(torch_model.children())
11    pre_layers = []
12    post_blocks = []
13
14    # for pre_layers
15    for layer in layers:
16        if type(layer) == nn.Sequential:
17            break
18        pre_layers.append(layer)
19    # for post_blocks
20    cnt = 0
21    for sequence in layers[5:]:
22        if type(sequence) == nn.Sequential:
23            for block in sequence:
24                post_blocks.append(block)
25            cnt += 1
26            if cnt > i:
27                break
28    # freezing (1)
29    for layer in pre_layers:
30        layer.requires_grad = False
31    # freezing (2)
32    for block in post_blocks:
33        block.requires_grad = False
34    final_layers = pre_layers + post_blocks
35    return nn.Sequential(*final_layers)
```

```
38 class ResNetLoss(nn.Module):
39     def __init__(self, i=5, device='cuda:0', resnet_model_name='resnet34'):
40         ''' [Logic]
41         1) extract the feature map from the 'i'-th residual block
42             for both hr_image (ground_truth) and sr_image (prediction).
43         2) compute and return mse_loss between features maps of hr_image & sr_image.
44         '''
45         super(ResNetLoss, self).__init__()
46         resnet = torch.hub.load('pytorch/vision:v0.10.0', resnet_model_name, pretrained=True).to(device).eval()
47         self.extractor = get_extractor_for_ResNet(resnet, i)
48         self.mse_loss = nn.MSELoss()
49
50     def forward(self, hr_img, sr_img):
51         hr_feature_map = self.extractor(hr_img)
52         sr_feature_map = self.extractor(sr_img)
53         vgg_loss = self.mse_loss(hr_feature_map, sr_feature_map)
54
55     return vgg_loss
```

SRResNet [Code]

SRGAN/1_SRResNet/SRResNet_train.py

61 lines (52 sloc) | 2.56 KB

Raw Blame

```
1 import sys
2 sys.path.append('..')
3
4 import datetime
5
6 from torch import optim
7 from torch.utils.data import DataLoader
8 from torch.utils.tensorboard import SummaryWriter
9
10 from Dataset.dataset import HR_to_HR_LR_PairDataset
11 from Model.Generator import Generator
12 from SRResNet_parser import Parser
13 from SRResNet_utils import *
14
15
16 if __name__ == '__main__':
17     # Parser
18     my_parser = Parser(mode='train')
19     my_args = my_parser.parser.parse_args()
20     print('*', my_args)
21
22     # Tag
23     cur_time = datetime.datetime.now().strftime('%y-%m-%d-%H-%M-%S')
24     tag_name = f'{my_args.tag}-{cur_time}' if my_args.tag else f'{cur_time}'
25     print('[SRResNet]')
26     print(f'{tag_name} experiment has been started.')
27
28     # Parameter
29     train_dir = my_args.train_dir
30     valid_dir = my_args.valid_dir
31     device = f'cuda:{my_args.gpu_index}' if my_args.gpu_index >= 0 else 'cpu'
32
33     # Loader (Train / Valid)
34     train_dataset = HR_to_HR_LR_PairDataset(train_dir, mode='train', crop_size=128)
35     valid_dataset = HR_to_HR_LR_PairDataset(valid_dir, mode='eval', crop_size=256)
36     train_loader = DataLoader(train_dataset, my_args.batch_size, shuffle=True, pin_memory=True)
37     valid_loader = DataLoader(valid_dataset, my_args.batch_size, shuffle=False, pin_memory=True)
```

```
39     # Train Generator
40     # if my_args.train_generator:
41     print('[Train a generator]')
42     generator = Generator(in_channels=3, scaling_factor=4, num_residual_blocks=16)
43     generator_best_model_state = generator.state_dict()
44     g_optimizer = optim.Adam(generator.parameters(), lr=my_args.g_lr, betas=(0.9, 0.999))
45     db_losses = {'train': [], 'valid': []}
46     tensorboard_writer = SummaryWriter(f'./log_dir/{tag_name}')
47     for epoch in range(my_args.epochs):
48         train_mse, valid_mse = \
49             train_and_validate_generator(generator, g_optimizer, epoch, device,
50                                         train_loader, valid_loader, my_args.loss_function, tensorboard_writer)
51         if db_losses['valid'] and \
52             valid_mse < min(db_losses['valid']): # store the best model so far.
53             generator_best_model_state = generator.state_dict()
54         db_losses['train'].append(train_mse)
55         db_losses['valid'].append(valid_mse)
56     torch.save(generator_best_model_state, f'SRResNet-generator-{tag_name}.pt')
57     print(f'SRResNet-generator-{tag_name}.pt is stored.')
58     write_log_for_SRResNet(db_losses, f'SRResNet-generator-{tag_name}-log.txt')
59     print(f'SRResNet-generator-{tag_name}-log.txt is stored.')
60     plot_log_for_SRResNet(db_losses, f'SRResNet-generator-{tag_name}-log.png')
61     print(f'SRResNet-generator-{tag_name}-log.png is stored.'
```

SRResNet [Code]

SRGAN/1_SRResNet/SRResNet_utils.py

```
30 def train_and_validate_generator(generator, g_optimizer, epoch, device,
31                                 train_loader, valid_loader, loss_function_name, tensorboard_writer):
32     generator.to(device)
33     loss_function = nn.MSELoss().to(device)
34     if loss_function_name == 'vgg_loss_19_5_4':
35         loss_function = VGGLoss(i=5, j=4, vgg_model_name='vgg19_bn')
36     elif loss_function_name == 'res_loss_18_4':
37         loss_function = ResNetLoss(i=4, resnet_model_name='resnet18')
38     elif loss_function_name == 'res_loss_34_3':
39         loss_function = ResNetLoss(i=3, resnet_model_name='resnet34')
40     elif loss_function_name == 'res_loss_34_5':
41         loss_function = ResNetLoss(i=5, resnet_model_name='resnet34')
42
43     # (1) train phase
44     train_mse = 0
45     generator.train()
46     train_tqdm_loader = tqdm.tqdm(train_loader)
47     for i, (lr_img, hr_img) in enumerate(train_tqdm_loader):
48         lr_img = lr_img.to(device)
49         hr_img = hr_img.to(device)
50
51         sr_img = generator(lr_img)
52         train_loss = loss_function(hr_img, sr_img)
53         train_mse += train_loss.item()
54
55         # optimization
56         g_optimizer.zero_grad()
57         train_loss.backward()
58         g_optimizer.step()
59
60         train_tqdm_loader.set_description(f'Train-SRResNet | Epoch: {epoch + 1} | Loss: {train_mse / (i + 1): .4f}')
61
62         # tensorboard log
63         train_mse /= len(train_tqdm_loader)
64         tensorboard_writer.add_scalar('Loss/train', train_mse, epoch + 1)
65         time.sleep(0.01)
66
67     # (2) validation phase
68     valid_mse = 0
69     generator.eval()
70     valid_tqdm_loader = tqdm.tqdm(valid_loader)
71     with torch.no_grad():
72         for i, (lr_img, hr_img) in enumerate(valid_tqdm_loader):
73             lr_img = lr_img.to(device)
74             hr_img = hr_img.to(device)
75
76             sr_img = generator(lr_img)
77             valid_loss = loss_function(hr_img, sr_img)
78             valid_mse += valid_loss.item()
79
80             valid_tqdm_loader.set_description(f'Valid-SRResNet | Epoch: {epoch + 1} | '
81                                              f'Loss: {valid_mse / (i + 1): .4f}')
82
83     valid_mse /= len(valid_tqdm_loader)
84     tensorboard_writer.add_scalar('Loss/valid', valid_mse, epoch + 1)
85
86     return train_mse, valid_mse
```

Training Details for SRGAN

- During training, the learning rate is 0.0001 and no learning rate scheduler is adopted.
- Two ADAM optimizers are used for both discriminator and generator.
- ADAM optimizer is used with betas=(0.9, 0.999).

- SRGAN models are trained for 100 / 200 epochs on DIV2K.

- We consider following content losses as follows:

- (1) MSE [+a]
 - (2) VGG Loss [Default]
 - (3) ResNet Loss [+a]

→ See the above slide for more details!

- Discriminator

- Training objective

$$\begin{aligned} & \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))] \\ & \Leftrightarrow \min_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [-\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [-\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))] \end{aligned}$$

(For PyTorch Implementation)

- Generator

↔ SRResNet

- Generator is initialized with pretrained weights of SRResNet (w/ MSE Loss).
 - Training objective

$$\begin{aligned} & \min_{\theta_G} \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))] \\ & \Leftrightarrow \min_{\theta_G} \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [-\log(D_{\theta_D}(G_{\theta_G}(I^{LR})))] \end{aligned}$$

(For PyTorch Implementation)

Training Details for SRGAN

- Discriminator

- Training objective [Code]

$$\begin{aligned} & \max_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))] \\ \Leftrightarrow & \min_{\theta_D} \mathbb{E}_{I^{HR} \sim p_{train}(I^{HR})} [-\log D_{\theta_D}(I^{HR})] + \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [-\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))] \end{aligned}$$

SRGAN/2_SRGAN/SRGAN_utils.py

```
55     discriminator.train()
56     generator.train()
57     train_tqdm_loader = tqdm.tqdm(train_loader)
58     for i, (lr_img, hr_img) in enumerate(train_tqdm_loader):
59         lr_img = lr_img.to(device)
60         hr_img = hr_img.to(device)
61
62         # (1-1) train a discriminator model.
63         d_optimizer.zero_grad()
64         sr_img = generator(lr_img).detach()      # detach() <- for not back-propagation
65         d_i_hr = discriminator(hr_img).flatten()
66         d_i_sr = discriminator(sr_img).flatten()
67         d_loss_real = bce_loss_function(d_i_hr, torch.ones_like(d_i_hr))    # Goal: D[I^{HR}] ~ 1 (for real image)
68         d_loss_fake = bce_loss_function(d_i_sr, torch.zeros_like(d_i_sr))   # Goal: D[G(I^{LR})] ~ 0 (for fake image)
69         d_loss = d_loss_real + d_loss_fake        # maximize E[log[D[I^{HR}]]] + E[log[(1 - D[G(I^{LR})])]]
70                                         # <=> maximize E[log[D[I^{HR}]]] - E[log[D[G(I^{LR})]]]
71                                         # for PyTorch implementation (PyTorch only has a minimizer!!!),
72                                         # <=> minimize - E[log[D[I^{HR}]]] + E[log[D[G(I^{LR})]]]
73         train_d_loss += d_loss.item()
74         d_loss.backward(retain_graph=True)
75         d_optimizer.step()
```

Binary Cross Entropy

$$L(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

$$\begin{aligned} \text{if } y = 1, \quad L &= -\log \hat{y} \\ \text{otherwise } (y = 0), \quad L &= -\log(1 - \hat{y}) \end{aligned}$$

Training Details for SRGAN

- Generator

↔ SRResNet

- Generator is initialized with pretrained weights of SRResNet (w/ MSE Loss).
- Training objective [Code]

$$\begin{aligned} & \min_{\theta_G} \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [\log(1 - D_{\theta_D}(G_{\theta_G}(I^{LR})))] \\ & \Leftrightarrow \min_{\theta_G} \mathbb{E}_{I^{LR} \sim p_G(I^{LR})} [-\log(D_{\theta_D}(G_{\theta_G}(I^{LR})))] \end{aligned}$$

Binary Cross Entropy

$$L(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

$$\begin{aligned} \text{if } y = 1, \quad L &= -\log \hat{y} \\ \text{otherwise } (y = 0), \quad L &= -\log(1 - \hat{y}) \end{aligned}$$

SRGAN/2_SRGAN/SRGAN_utils.py

```
77      # (1-2) train a generator model.
78      g_optimizer.zero_grad()
79      hr_img = hr_img.detach()           # detach() <- for not back-propagation
80      sr_img = generator(lr_img)
81      d_i_sr = d_i_sr.detach()
82      content_loss = content_loss_function(hr_img, sr_img)
83      adversarial_loss = bce_loss_function(d_i_sr, torch.ones_like(d_i_sr)) # minimize E[log[(1 - D[G(I^{LR})])]]
84                                         # <=> minimize - E[log[D[G(I^{LR})]]]
85      g_loss = 1 * content_loss + 1e-3 * adversarial_loss
86
87      train_g_loss += g_loss.item()
88      g_loss.backward()
89      g_optimizer.step()
```

SRGAN [Code]

SRGAN/2_SRGAN/SRGAN_train.py

```
17 if __name__ == '__main__':
18     # Parser
19     my_parser = Parser(mode='train')
20     my_args = my_parser.parser.parse_args()
21
22     # Tag
23     cur_time = datetime.datetime.now().strftime('%y-%m-%d-%H-%M-%S')
24     tag_name = f'{my_args.tag}-{cur_time}' if my_args.tag else f'{cur_time}'
25     print('[SRGAN]')
26     print(f'{tag_name} experiment has been started.')
27
28     # Parameter
29     train_dir = my_args.train_dir
30     valid_dir = my_args.valid_dir
31     device = f'cuda:{my_args.gpu_index}' if my_args.gpu_index >= 0 else 'cpu'
32
33     # Loader (Train / Valid)
34     train_dataset = HR_to_HR_LR_PairDataset(train_dir, mode='train', crop_size=128)
35     valid_dataset = HR_to_HR_LR_PairDataset(valid_dir, mode='eval', crop_size=256)
36     train_loader = DataLoader(train_dataset, my_args.batch_size, shuffle=True, pin_memory=True)
37     valid_loader = DataLoader(valid_dataset, my_args.batch_size, shuffle=False, pin_memory=True)
38
39     discriminator = Discriminator(in_channels=3)
40     discriminator_best_psnr_model_state = discriminator.state_dict()
41     discriminator_best_ssim_model_state = discriminator.state_dict()
42     d_optimizer = optim.Adam(discriminator.parameters(), lr=my_args.d_lr, betas=(0.9, 0.999))
43
44     generator = Generator(in_channels=3, scaling_factor=4, num_residual_blocks=16)
45     generator.load_state_dict(torch.load(my_args.generator_pt_path))
46     generator_best_psnr_model_state = generator.state_dict()
47     generator_best_ssim_model_state = generator.state_dict()
48     g_optimizer = optim.Adam(generator.parameters(), lr=my_args.g_lr, betas=(0.9, 0.999))
49
50     db_valid_PSNR_SSIM = {'psnr': [], 'ssim': []}
51     tensorboard_writer = SummaryWriter(f'./log_dir/{tag_name}')
53
54         for epoch in range(my_args.epochs):
55             valid_psnr, valid_ssim = \
56                 train_and_validate_SRGAN(generator, g_optimizer, discriminator, d_optimizer, epoch, device,
57                                         train_loader, valid_loader, my_args.content_loss_function, tensorboard_writer)
58
59             if db_valid_PSNR_SSIM['psnr'] and \
60                 valid_ssim > max(db_valid_PSNR_SSIM['psnr']): # store the best model so far.
61                 generator_best_psnr_model_state = generator.state_dict()
62                 discriminator_best_psnr_model_state = discriminator.state_dict()
63
64             if db_valid_PSNR_SSIM['ssim'] and \
65                 valid_ssim > max(db_valid_PSNR_SSIM['ssim']): # store the best model so far.
66                 generator_best_ssim_model_state = generator.state_dict()
67                 discriminator_best_ssim_model_state = discriminator.state_dict()
68
69             db_valid_PSNR_SSIM['psnr'].append(valid_psnr)
70             db_valid_PSNR_SSIM['ssim'].append(valid_ssim)
71
72             generator_last_model_state = generator.state_dict()
73             discriminator_last_model_state = discriminator.state_dict()
74
75             torch.save(generator_best_psnr_model_state, f'SRGAN-generator-best-psnr-{tag_name}.pt')
76             torch.save(generator_best_ssim_model_state, f'SRGAN-generator-best-ssim-{tag_name}.pt')
77             torch.save(generator_last_model_state, f'SRGAN-generator-last-{tag_name}.pt')
78
79             torch.save(discriminator_best_psnr_model_state, f'SRGAN-discriminator-best-psnr-{tag_name}.pt')
80             torch.save(discriminator_best_ssim_model_state, f'SRGAN-discriminator-best-ssim-{tag_name}.pt')
81             torch.save(discriminator_last_model_state, f'SRGAN-discriminator-last-{tag_name}.pt')
82
83             print(f'SRGAN-generator-best-psnr-{tag_name}.pt is stored.')
84             print(f'SRGAN-generator-best-ssim-{tag_name}.pt is stored.')
85             print(f'SRGAN-generator-last-{tag_name}.pt is stored.')
86
87             print(f'SRGAN-discriminator-best-psnr-{tag_name}.pt is stored.')
88             print(f'SRGAN-discriminator-best-ssim-{tag_name}.pt is stored.')
89             print(f'SRGAN-discriminator-last-{tag_name}.pt is stored.')
90
91             write_log_for_SRGAN(db_valid_PSNR_SSIM, f'SRGAN-{tag_name}-log.txt')
92             print(f'SRGAN-{tag_name}-log.txt is stored.')
93
94             plot_log_for_SRGAN(db_valid_PSNR_SSIM, f'SRGAN-{tag_name}-log.png')
95             print(f'SRGAN-{tag_name}-log.png is stored.'
```

Evaluation Metric (1)

- PSNR (Peak Signal-to-Noise Ratio)

- PSNR is formulated as follows:

$$PSNR(I_a, I_b) = 10 \log \frac{s^2}{MSE(I_a, I_b)}.$$



This work sets s to 1.
(\because Normalized image tensors are used.)

- I_a and I_b denote the pair of images of same dimension.

- The dimension of the image can be represented as (C, H, W) .

- C : Number of Channels
 - H : Height
 - W : Width

- s denotes the maximum possible pixel value of the image.

- $MSE(I_a, I_b) = \sum_{i=1}^C \sum_{j=1}^H \sum_{k=1}^W |I_a(i, j, k) - I_b(i, j, k)|^2$.

Reference:

[8] https://en.wikipedia.org/wiki/Peak_signal-to-noise_ratio

[9] <https://cvnote.ddlee.cc/2019/09/12/psnr-ssim-python>

Evaluation Metric (2)

- SSIM (Structural Similarity Index Map)

- SSIM is formulated as follows:

$$SSIM(I_a, I_b) = [l(I_a, I_b)^\alpha \cdot c(I_a, I_b)^\beta \cdot s(I_a, I_b)^\gamma].$$

- I_a and I_b denote the pair of images of same dimension.

- The dimension of the image can be represented as (C, H, W) .

- C : Number of Channels
 - H : Height
 - W : Width

- Luminance, $l(x, y) = \frac{2\mu_x\mu_y + c_1}{\mu_x^2 + \mu_y^2 + c_1}$

- μ_i : Average of i .

- Contrast, $c(x, y) = \frac{2\sigma_x\sigma_y + c_2}{\sigma_x^2 + \sigma_y^2 + c_2}$

- σ_i^2 : Variance of i .

- Structure, $s(x, y) = \frac{\sigma_{xy} + c_3}{\sigma_x\sigma_y + c_3}$

- σ_{xy} : Covariance of x and y

- In addition to above definitions, this work sets c_3 to $\frac{c_2}{2}$ and α, β, γ to 1,

- $SSIM(I_a, I_b) = \frac{(2\mu_x\mu_y + c_1)(2\sigma_{xy} + c_2)}{(\mu_x^2 + \mu_y^2 + c_1)(\sigma_x^2 + \sigma_y^2 + c_2)}$.



In this work, c_1 and c_2 are set to 0.01^2 and 0.03^2 , respectively as in [4].

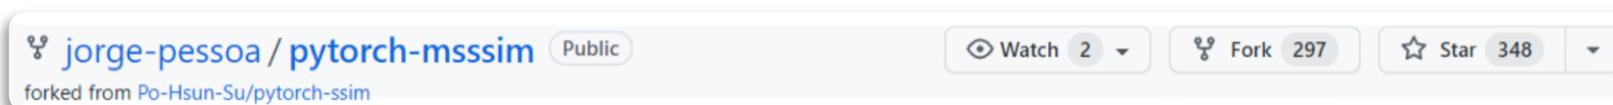
Reference:

[9] <https://cvnote.ddlee.cc/2019/09/12/psnr-ssim-python>

[10] https://en.wikipedia.org/wiki/Structural_similarity

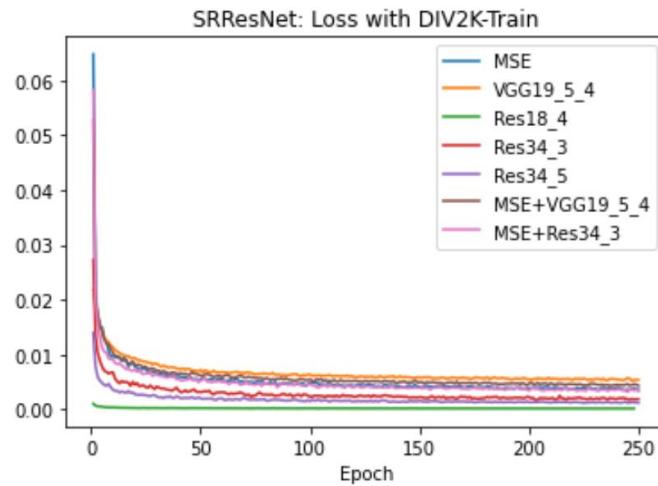
Experimental Setup

- OS / CPU: Linux Ubuntu 20.04 (LTS) 64bit, AMD Ryzen 5 3700x
- GPU: Nvidia RTX 3090 (24GB)
- CUDA: v11.2
- Python 3.7
 - PyTorch: Torch 1.10.0, Torchvision 0.11.1
 - Glob2
 - Tqdm
 - Tensorboard
 - Matplotlib
 - Pytorch-msssim
 - For **SSIM** → Rather than implementing SSIM myself, I used Jorge-pessoa's code.
 - Link: <https://github.com/jorge-pessoa/pytorch-msssim>
 - We can simply install and use Pytorch-msssim via pip installer (**pip install pytorch-msssim**).

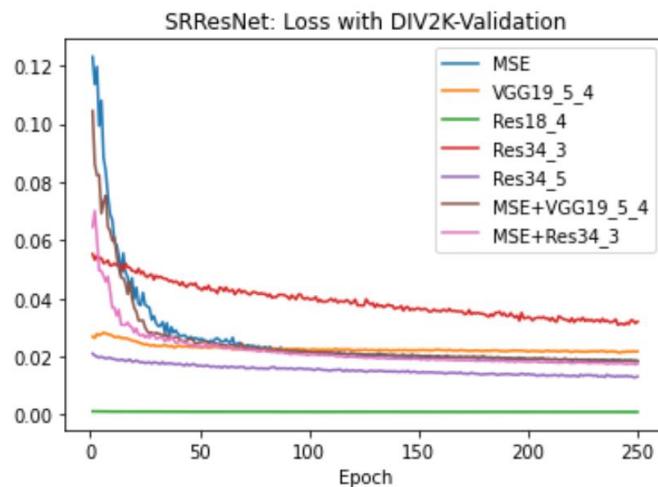
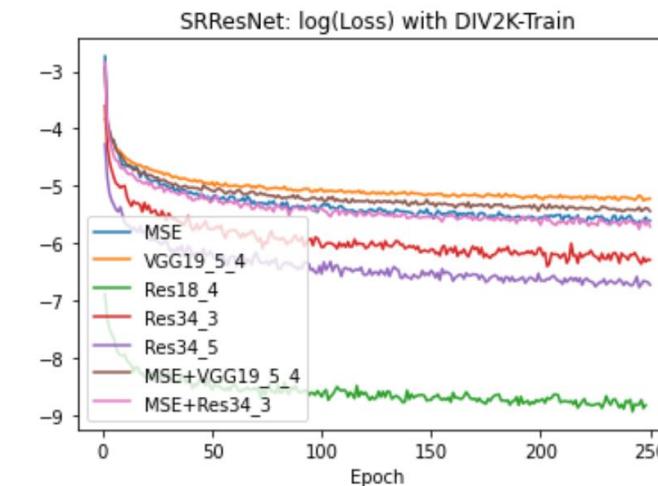


Experimental Results (1) SRResNet

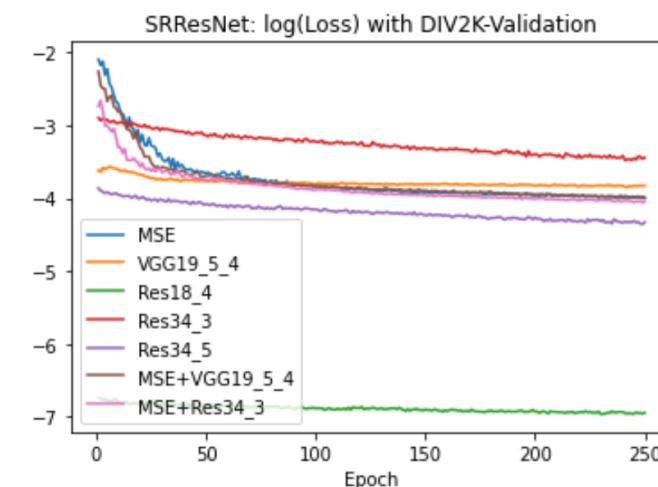
- Loss per each epoch



$\log(\cdot)$

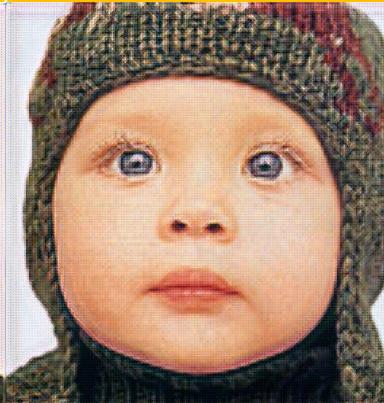


$\log(\cdot)$



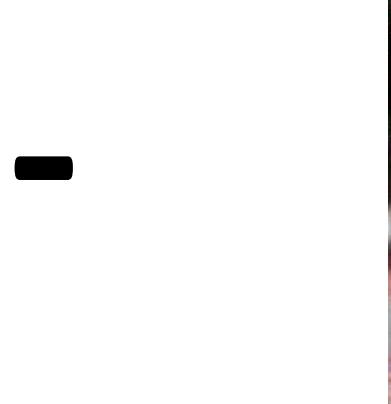
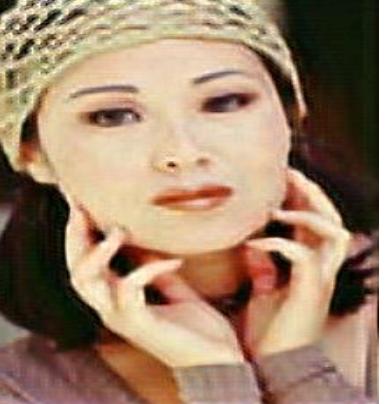
Experimental Results (1) SRResNet

- Visual Results on Set5 - Baby

GT	SRResNet-				
Set5/GTmod12/baby.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
				Bicubic	
0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3	0.7 * MSE + 0.3 * Res34_5	 Set5/LRbicx4/baby.png	
					

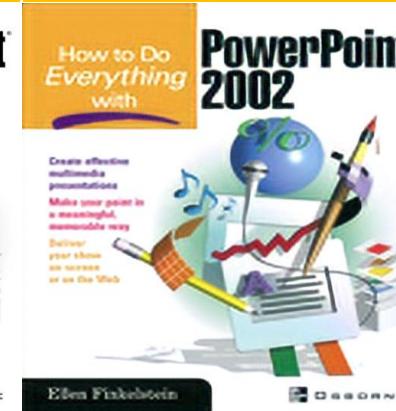
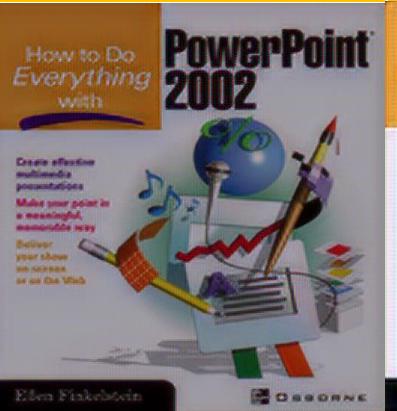
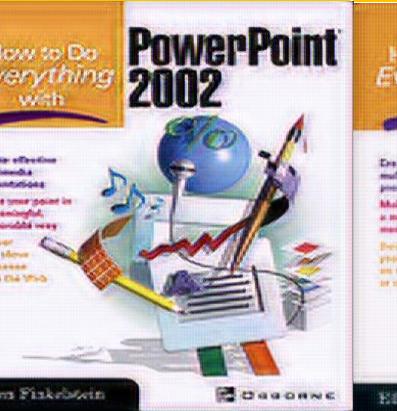
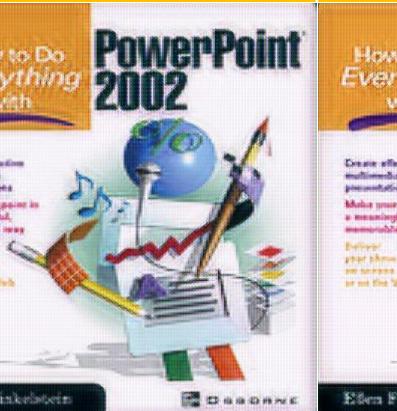
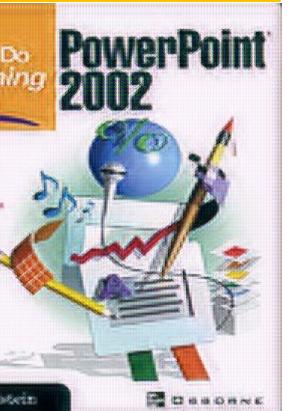
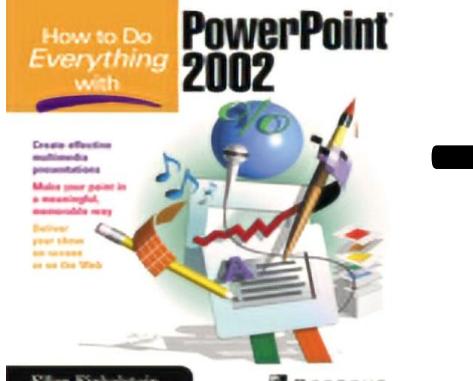
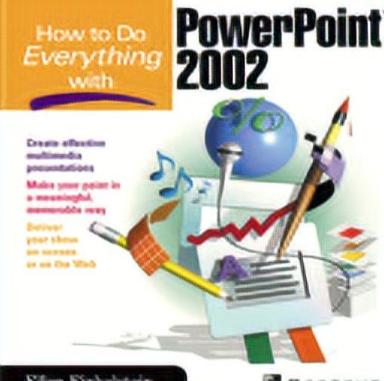
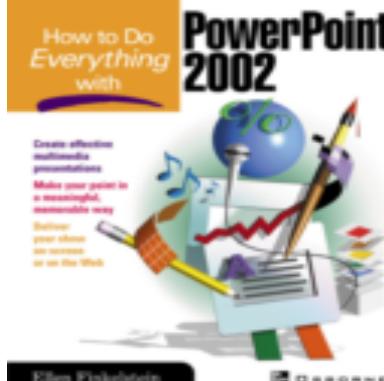
Experimental Results (1) SRResNet

- Visual Results on Set5 - Woman

GT	SRResNet-				
Set5/GTmod12/woman.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
				Bicubic	
0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3	0.7 * MSE + 0.3 * Res34_5	Set5/LRbicx4/woman.png	
					

Experimental Results (1) SRResNet

- Visual Results on Set14 - Ppt3

GT	SRResNet-					
Set14/GTmod12/ppt3.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	
						
Ellen Finkelstein	OSBORNE	Ellen Finkelstein	OSBORNE	Ellen Finkelstein	OSBORNE	Ellen Finkelstein
0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3	0.7 * MSE + 0.3 * Res34_5			Bicubic
						
Ellen Finkelstein	OSBORNE	Ellen Finkelstein	OSBORNE	Ellen Finkelstein	OSBORNE	

Experimental Results (1) SRResNet

- PSNR / SSIM on Set5 and Set14
 - Mean PSNR and Mean SSIM

SRResNet-									
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3	0.7 * MSE + 0.3 * Res34_5
PSNR	23.13	14.68	16.79	14.90	14.58	23.42	-	23.34	-
SSIM	0.9348	0.8186	0.8017	0.8287	0.8184	0.9379	-	0.9413	-

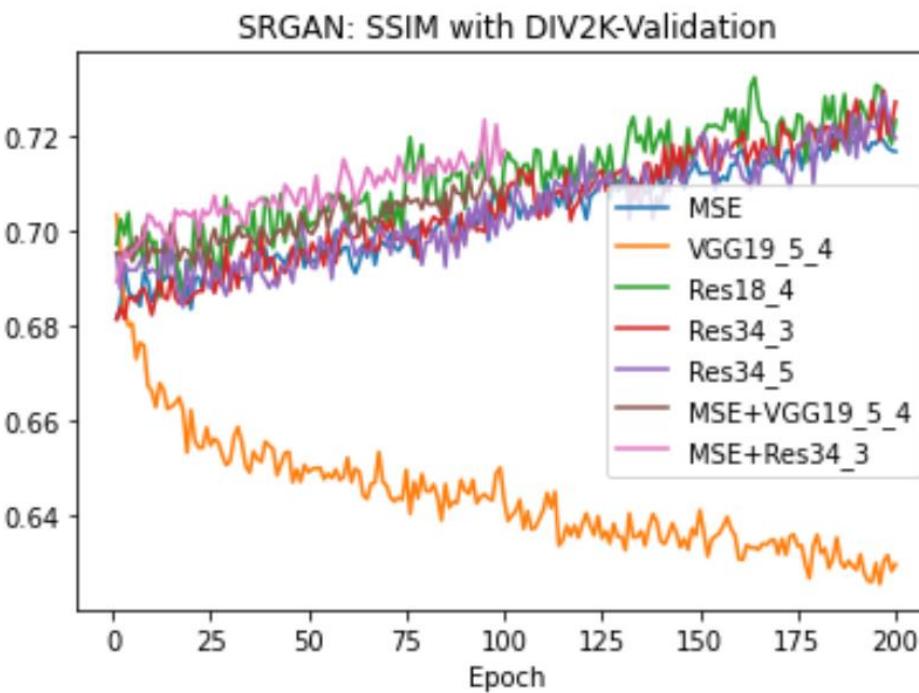
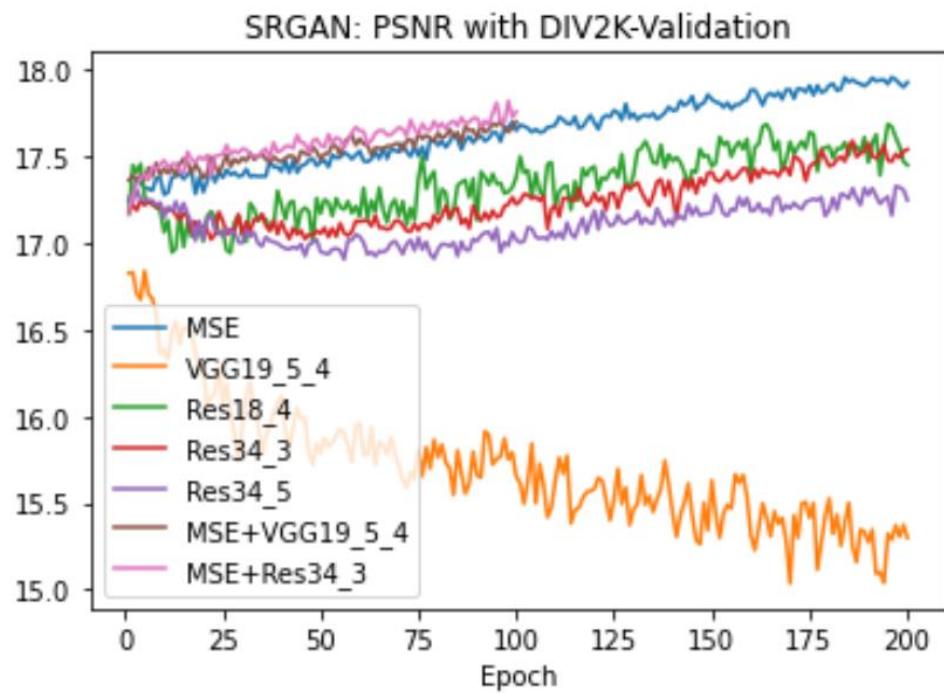
The diagram shows three curved arrows originating from the highlighted cells in the Set5 table and pointing to the corresponding cells in the Set14 table. One arrow points from the 'VGG19_5_4' cell in Set5 to the 'VGG19_5_4' cell in Set14. Another arrow points from the '0.7 * MSE + 0.3 * VGG19_5_4' cell in Set5 to the same cell in Set14. A third arrow points from the '0.7 * MSE + 0.3 * Res34_3' cell in Set5 to the same cell in Set14.

SRResNet-									
Set14	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3	0.7 * MSE + 0.3 * Res34_5
PSNR	21.73	14.00	16.29	14.67	14.37	22.16	-	21.99	-
SSIM	0.9090	0.8064	0.7812	0.8104	0.7989	0.9112	-	0.9164	-

The diagram shows three curved arrows originating from the highlighted cells in the Set14 table and pointing to the corresponding cells in the Set5 table. One arrow points from the 'VGG19_5_4' cell in Set14 to the 'VGG19_5_4' cell in Set5. Another arrow points from the '0.7 * MSE + 0.3 * VGG19_5_4' cell in Set14 to the same cell in Set5. A third arrow points from the '0.7 * MSE + 0.3 * Res34_3' cell in Set14 to the same cell in Set5.

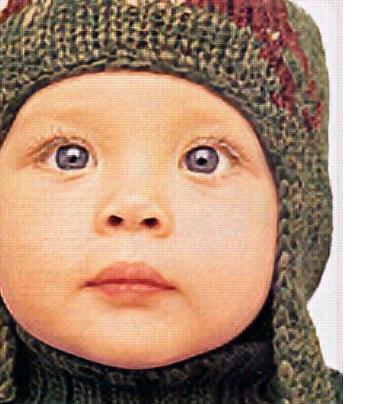
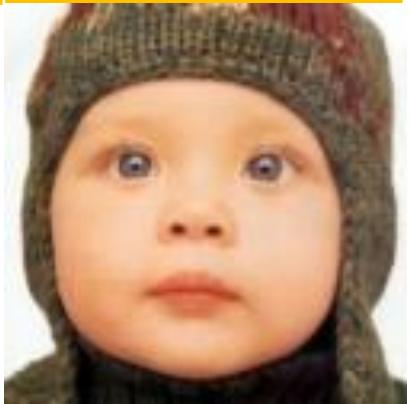
Experimental Results (2) SRGAN

- Validation PSNR / SSIM per each epoch



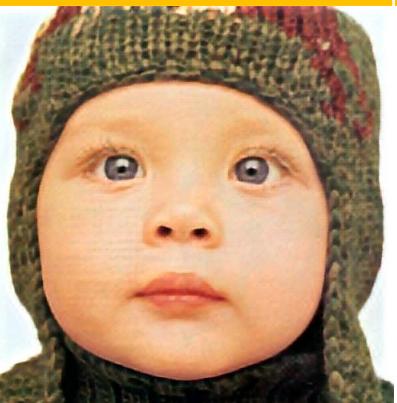
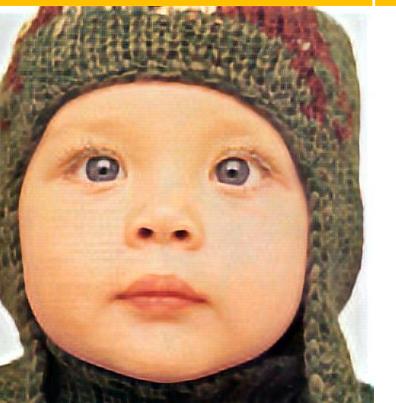
Experimental Results (2) SRGAN

- Visual Results on Set5 - Baby

GT	SRGAN-	(200 epochs)			
Set5/GTmod12/baby.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
				Bicubic	
0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5	Set5/LRbicx4/baby.png	

Experimental Results (2) SRGAN

- Visual Results on Set5 - Baby

GT	SRGAN-	(100 epochs)			
Set5/GTmod12/baby.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
				Bicubic	
0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5	Set5/LRbicx4/baby.png	
					

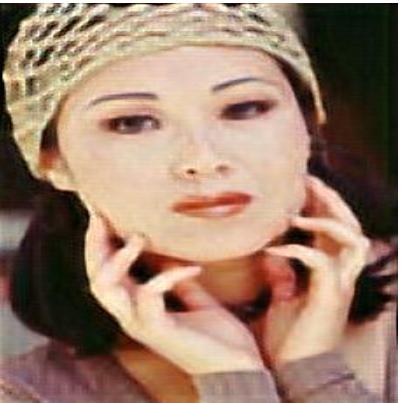
Experimental Results (2) SRGAN

- Visual Results on Set5 - Woman

GT	SRGAN-	(200 epochs)			
Set5/GTmod12/woman.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
				Bicubic	
0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5	Set5/LRbicx4/woman.png	

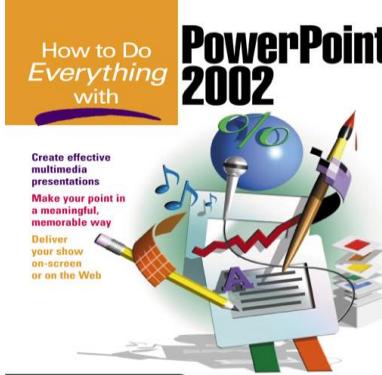
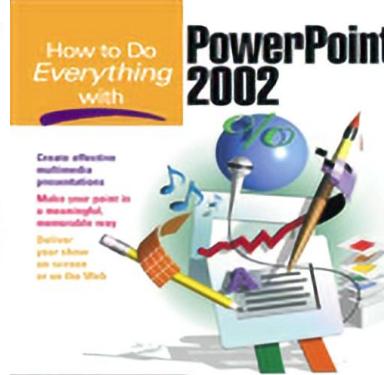
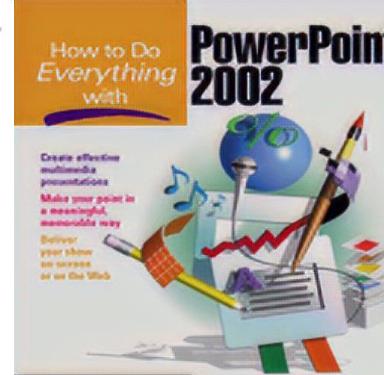
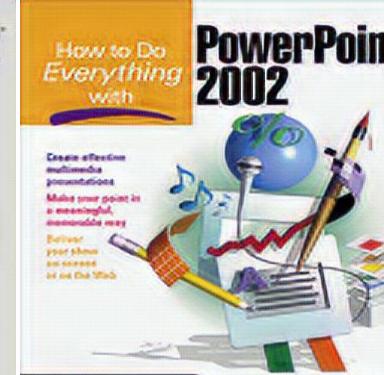
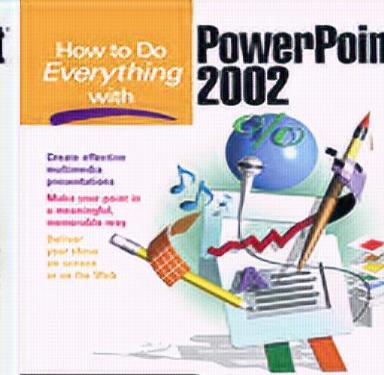
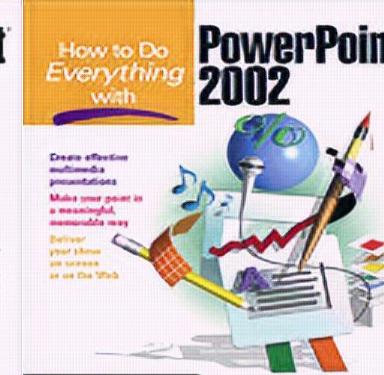
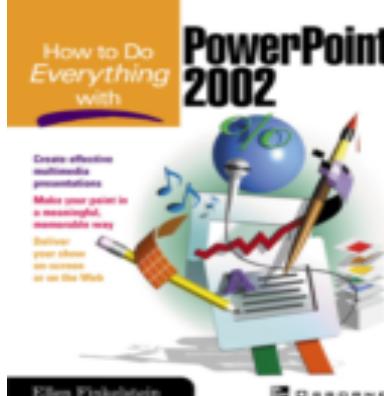
Experimental Results (2) SRGAN

- Visual Results on Set5 - Woman

GT	SRGAN-	(100 epochs)			
Set5/GTmod12/woman.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
				Bicubic	
0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5	Set5/LRbicx4/woman.png	
					

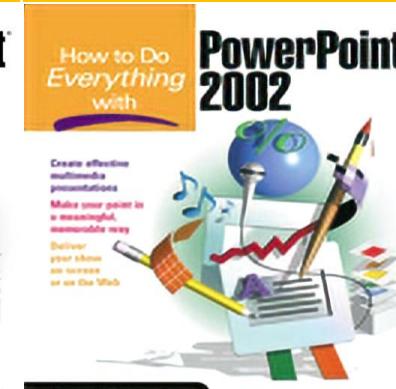
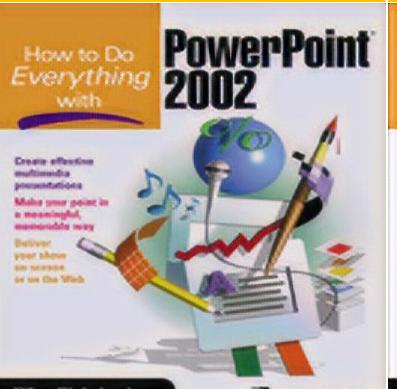
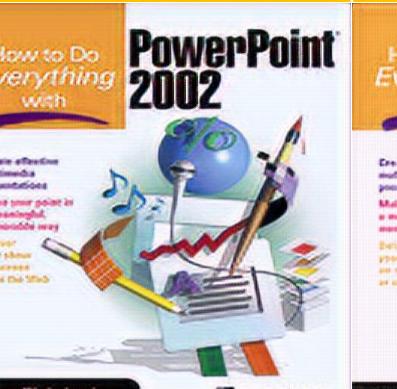
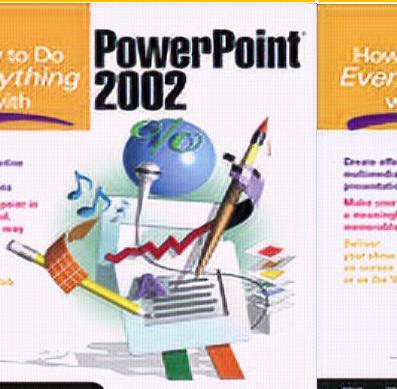
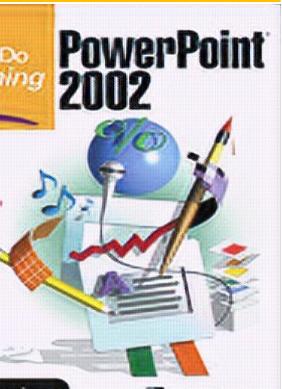
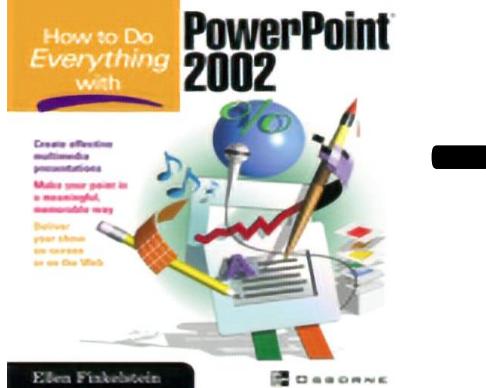
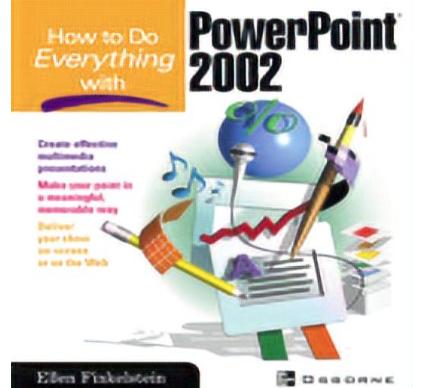
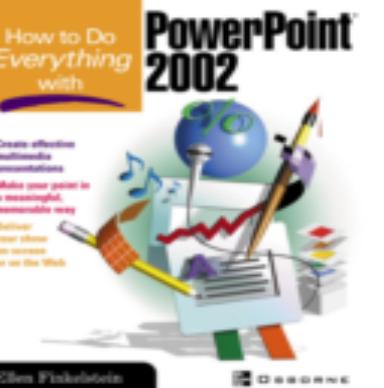
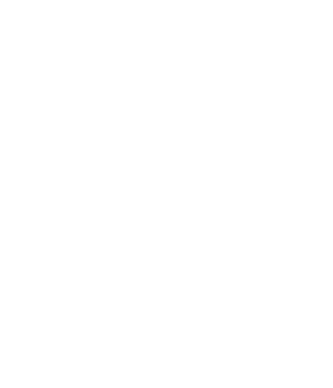
Experimental Results (2) SRGAN

- Visual Results on Set14 - Ppt3

GT	SRGAN-	(200 epochs)				
Set14/GTmod12/ppt3.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	
						
Ellen Finkelstein	OSBORNE	Ellen Finkelstein	OSBORNE	Ellen Finkelstein	OSBORNE	Ellen Finkelstein
0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5	Set5/LRbicx4/ppt3.png		

Experimental Results (2) SRGAN

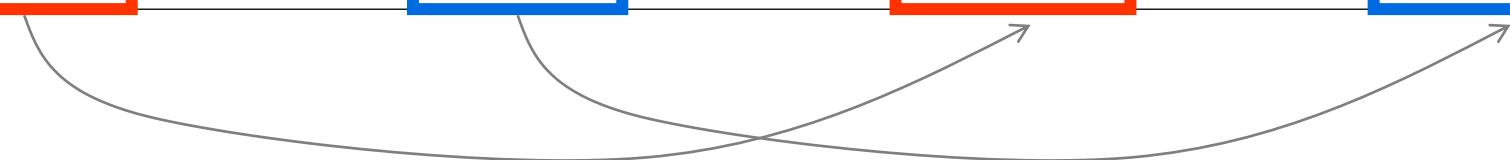
- Visual Results on Set14 - Ppt3

GT	SRGAN-	(100 epochs)			
Set14/GTmod12/ppt3.png	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5
					
Ellen Finkelstein	Ellen Finkelstein	Ellen Finkelstein	Ellen Finkelstein	Ellen Finkelstein	Ellen Finkelstein
OSBORNE	OSBORNE	OSBORNE	OSBORNE	OSBORNE	OSBORNE
0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5	Set5/LRbicx4/ppt3.png	Bicubic
					

Experimental Results (2) SRGAN

- PSNR / SSIM on Set5 and Set14
 - Mean PSNR and Mean SSIM

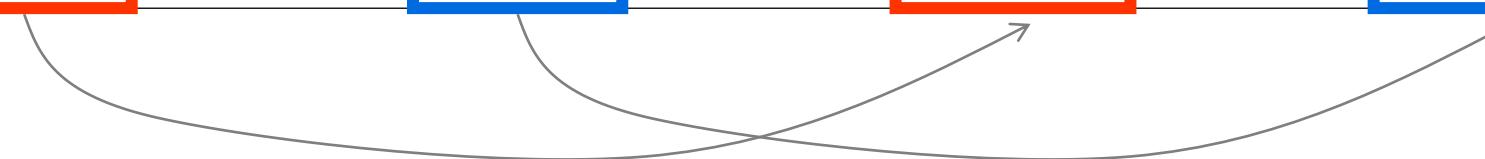
SRGAN- 200 epochs									
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5
PSNR	24.17	20.12	20.21	18.88	18.75	-	-	-	-
SSIM	0.9470	0.8926	0.8917	0.8754	0.8706	-	-	-	-
SRGAN- 100 epochs									
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5
PSNR	23.07	20.96	20.44	19.62	19.24	23.48	-	23.29	-
SSIM	0.9364	0.9073	0.8717	0.8891	0.8769	0.9426	-	0.9375	-



Experimental Results (2) SRGAN

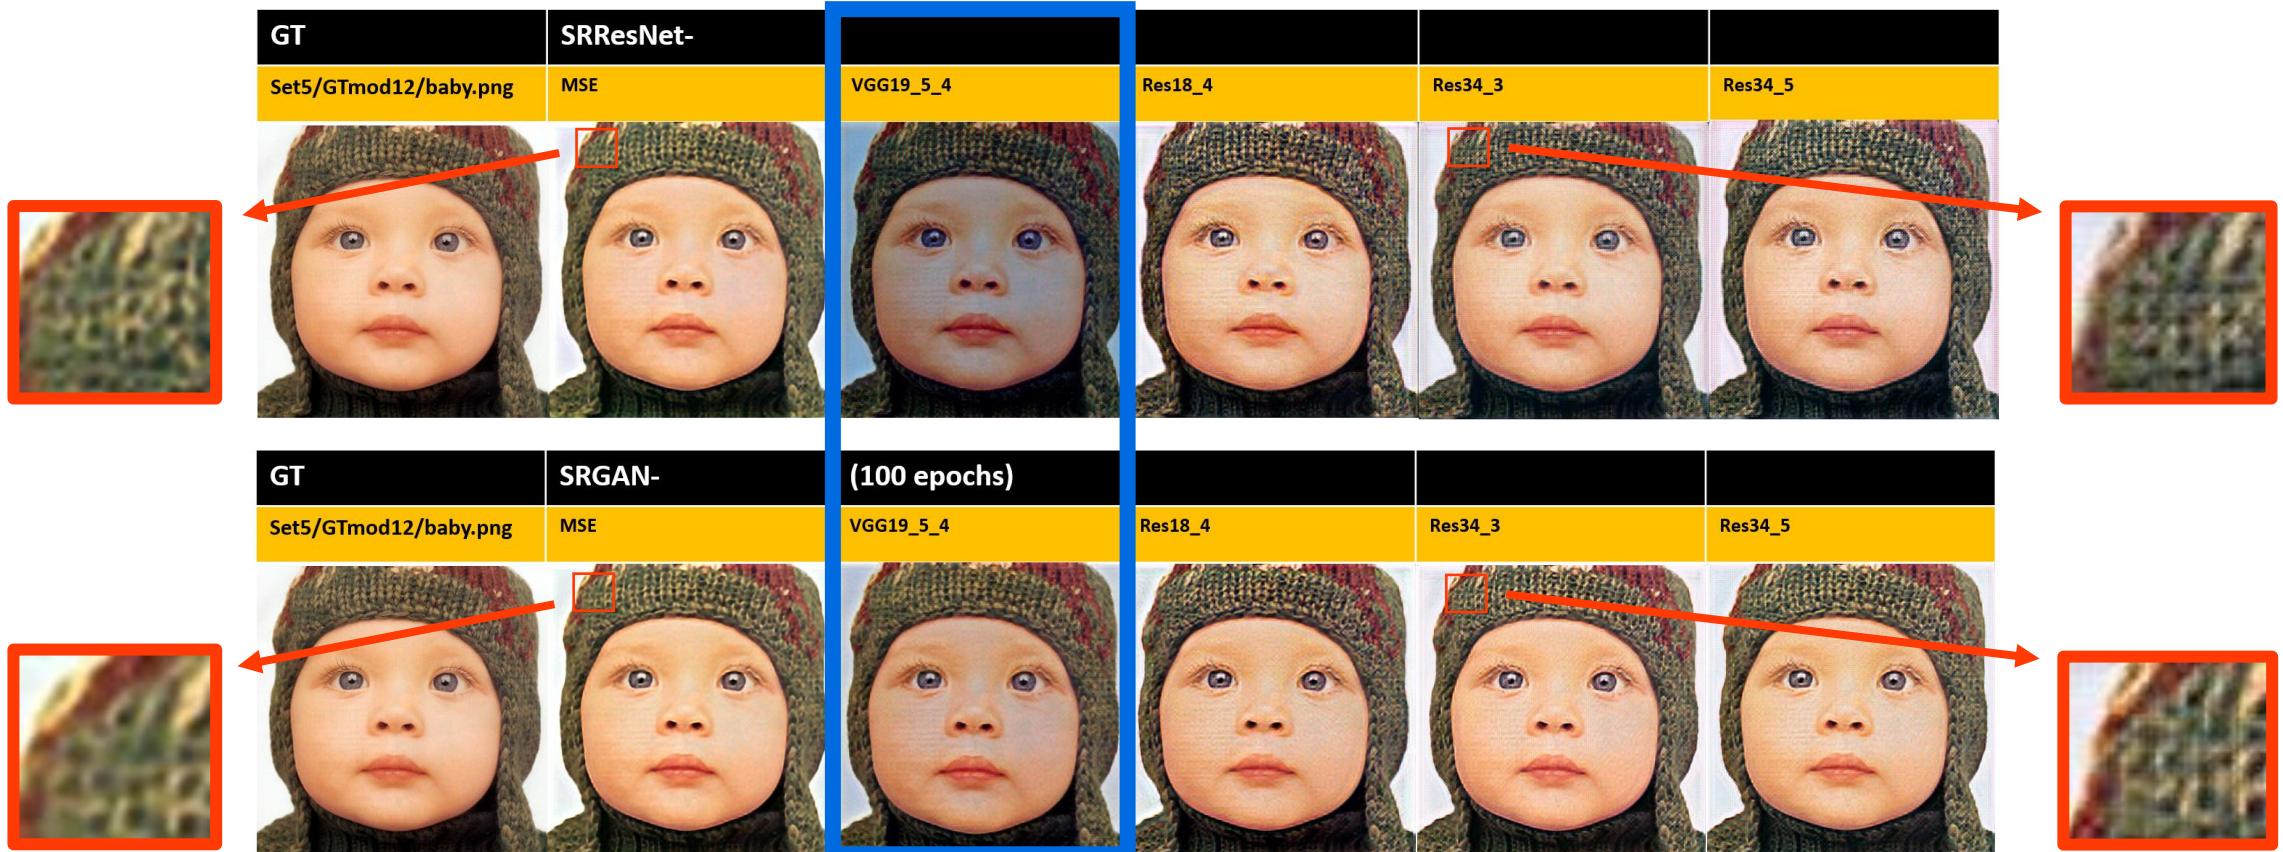
- PSNR / SSIM on Set5 and Set14
 - Mean PSNR and Mean SSIM

SRGAN- 200 epochs									
Set14	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5
PSNR	22.50	19.40	19.29	18.03	18.02	-	-	-	-
SSIM	0.9203	0.8783	0.8462	0.8465	0.8449	-	-	-	-
SRGAN- 100 epochs									
Set14	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5
PSNR	21.73	20.33	19.59	18.88	18.42	22.21	-	21.96	-
SSIM	0.9106	0.8920	0.8485	0.8561	0.8494	0.9197	-	0.9120	-



Conclusion

A. **Finer texture details can be recovered with VGG-based and ResNet-based content loss function.**



[Note] 1) Content loss function has more impact on the super resolution task than GAN!

2) Of course, GAN also boosts the performance of super resolution (See the 'blue' column □.).!

Conclusion

B. **Generative Adversarial Network** boosts the super resolution performance in terms of **PSNR & SSIM** in this work.

SRResNet-									
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3	0.7 * MSE + 0.3 * Res34_5
PSNR	23.13	14.68	16.79	14.90	14.58	23.42	-	23.34	-
SSIM	0.9348	0.8186	0.8017	0.8287	0.8184	0.9379	-	0.9413	-

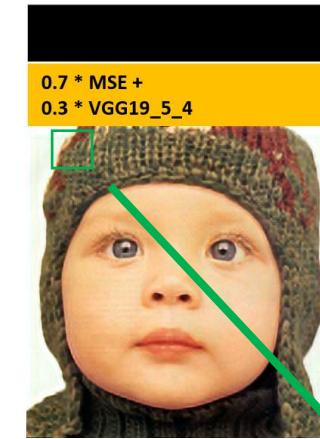
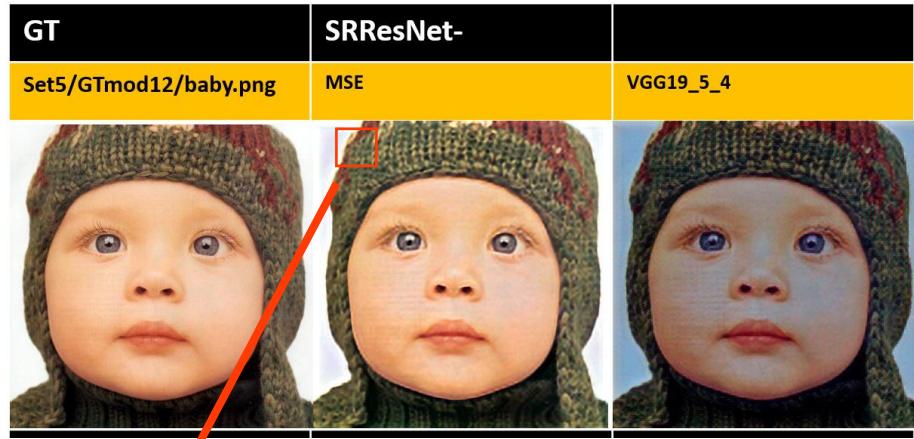


SRGAN- 200 epochs									
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5
PSNR	24.17	20.12	20.21	18.88	18.75	-	-	-	-
SSIM	0.9470	0.8926	0.8917	0.8754	0.8706	-	-	-	-

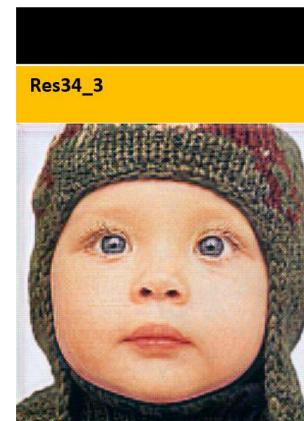
SRGAN- 100 epochs									
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3	0.6 * MSE + 0.4 * Res34_5
PSNR	23.07	20.96	20.44	19.62	19.24	23.48	-	23.29	-
SSIM	0.9364	0.9073	0.8717	0.8891	0.8769	0.9426	-	0.9375	-

Conclusion

C. Hybrid of ‘Mean Squared Error Loss’ and ‘Content Loss’ improves the quality of super resolution results.



Artifact ↓
Texture Detail ↑



Conclusion

C. Hybrid of 'Mean Squared Error Loss' and 'Content Loss' improves the quality of super resolution results.



Conclusion

C. Hybrid of ‘Mean Squared Error Loss’ and ‘Content Loss’ improves the quality of super resolution results.

SRResNet-								
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3
PSNR	23.13	14.68	16.79	14.90	14.58	23.42	-	23.34
SSIM	0.9348	0.8186	0.8017	0.8287	0.8184	0.9379	-	0.9413

SRResNet-								
Set14	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.7 * MSE + 0.3 * VGG19_5_4	0.7 * MSE + 0.3 * Res18_4	0.7 * MSE + 0.3 * Res34_3
PSNR	21.73	14.00	16.29	14.67	14.37	22.16	-	21.99
SSIM	0.9090	0.8064	0.7812	0.8104	0.7989	0.9112	-	0.9164

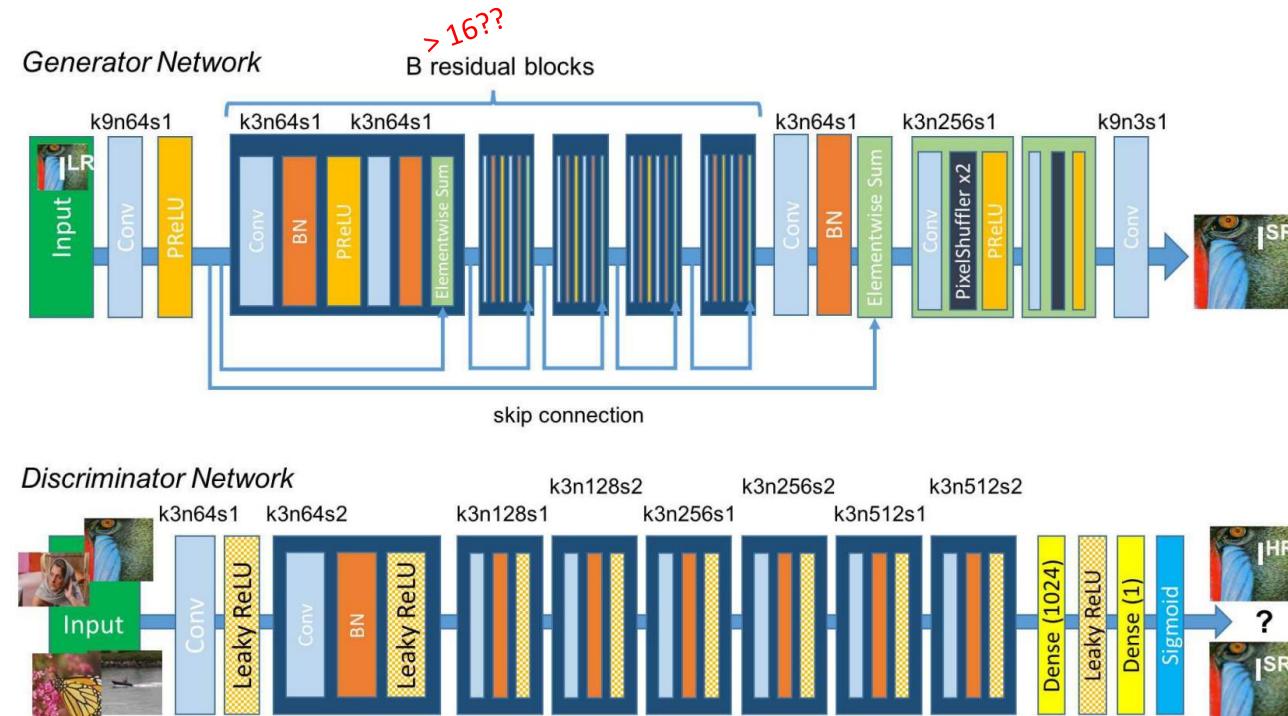
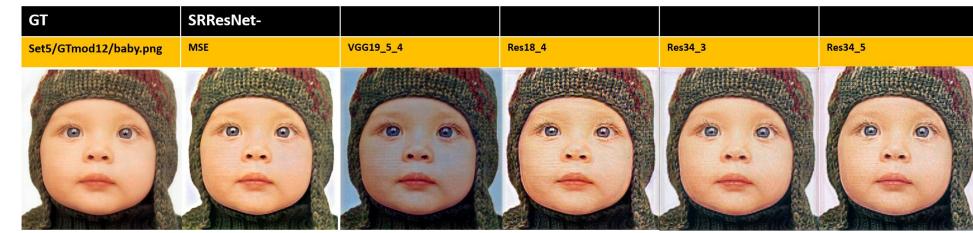
SRGAN- 100 epochs								
Set5	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3
PSNR	23.07	20.96	20.44	19.62	19.24	23.48	-	23.29
SSIM	0.9364	0.9073	0.8717	0.8891	0.8769	0.9426	-	0.9375

SRGAN- 100 epochs								
Set14	MSE	VGG19_5_4	Res18_4	Res34_3	Res34_5	0.6 * MSE + 0.4 * VGG19_5_4	0.6 * MSE + 0.4 * Res18_4	0.6 * MSE + 0.4 * Res34_3
PSNR	21.73	20.33	19.59	18.88	18.42	22.21	-	21.96
SSIM	0.9106	0.8920	0.8485	0.8561	0.8494	0.9197	-	0.9120

Both PSNR and SSIM are also improved!!!

Discussion

- In the future work,
 - We will investigate why VGG content loss results in color-desaturation of super-resolution results, compared to ResNet-based loss.
 - We will design more various content loss with other CNNs.
 - We will increase the number (B) of residual blocks in the generator in SRGAN.
 - $B > 16$



References

- [1] Ledig, Christian, et al. "Photo-realistic single image super-resolution using a generative adversarial network." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2017.
- [2] <https://kh-kim.gitbook.io/natural-language-processing-with-pytorch/00-cover-11/01-intro>
- [3] <https://www.slideshare.net/NaverEngineering/1-gangenerative-adversarial-network>
- [4] He, Kaiming, et al. "Deep residual learning for image recognition." Proceedings of the IEEE conference on computer vision and pattern recognition. 2016.
- [5] <https://paperswithcode.com/dataset/div2k>
- [6] <https://paperswithcode.com/dataset/set5>
- [7] <https://paperswithcode.com/dataset/set11>
- [8] https://en.wikipedia.org/wiki/Peak_signal-to-noise_ratio
- [9] <https://cvnote.ddlee.cc/2019/09/12/psnr-ssim-python>
- [10] https://en.wikipedia.org/wiki/Structural_similarity